

Long-context inference throughput (Llama-2-7B, 8×A100, NF4 + ring SP + YaRN)

