When Speculative Decoding Stops Paying Off at Long Context: A Diagnostic Study with RASD

Aman Kesarwani*[1]

[1]University of California, Berkeley (alumnus)

October 13, 2026

Abstract

Speculative decoding's short-context speedups do not survive million-token inference, and the mechanism is not what the technique's own literature would predict. We report three scientific findings and one systems measurement, all on 8×A100 80 GB SXM4 in FP4 weights with an NF4 KV cache. *First*, on `meta-llama/Llama-3.1-8B` — a model that ships the rotary-scaling dictionary the intervention modifies — stretching the rotary basis at 128 k moves paired perplexity by +0.05% (factor 16) and +1.56% (factor 32) while speculative acceptance falls from 0.7272 to 0.5336/0.5006, a 27–31% relative loss. An intervention that is nearly free in likelihood is expensive in acceptance, so acceptance is not a proxy for perplexity and must be measured. *Second*, greedy decoding at long context degenerates into repetition — 9 of 15 long-context rows loop against 0 of 26 sampled rows at T=0.7 — and a looping row is trivially accepted by a greedy draft, so a greedy acceptance number measures the degeneracy rather than the draft's agreement with the target. *Third*, the native-1M target stays coherent at 1 M: within 9.5% of its own 128 k perplexity on the exactly aligned book. *Fourth*, in the systems section we localise the decode bottleneck of the ring-attention engine we built: rotating each layer's key/value shard around 8 ranks costs 256 rounds per decode token, and a bare-process benchmark of the ring primitive shows each round costs a fixed ∼10.3 ms *regardless of payload* (4 KB and 266 MB are within 0.4% of each other). The decode step was therefore latency per round, not bandwidth. Removing the rotation with split-KV — prior work, Meta's pass-Q/pass-KV and vLLM's Decode Context Parallelism — is worth 2.8–5.3× in decode time per token, and speculative decoding still pays 2.0–2.5× on top of it. The measurement also withdraws the reviewed version's two systems claims, that ring communication was ≤1.2% of wall time and that the speedup regressed at 512 k/1 M.

Keywords: speculative decoding, ring attention, long-context LLM inference, KV cache quantization, RoPE scaling, decode-side context parallelism, distributed inference

1 Introduction

Long-context language model inference has emerged as the rate-limiting step for several practical workloads: multi-document question answering over book-length corpora, agentic workflows that accumulate hours of tool-call traces, retrieval-free in-context learning at the scale of a small library, and analysis of complete code repositories. Frontier closed-source systems already advertise context windows in the millions of tokens. Reproducing that capability on commodity, open hardware — 8 GPUs on a single node, the configuration available to most academic and small-industry groups — requires solving two compounding bottlenecks simultaneously.

*amankesarwani@berkeley.edu

Bottleneck 1 (memory). The key-value (KV) cache grows linearly with sequence length. For Llama-2-7B in bf16, a 1 million-token KV cache occupies approximately 512 GiB — far exceeding the 80 GiB capacity of a single A100 — and the attention matrix for prefill scales quadratically. Ring Attention [Liu et al., 2023] resolves this by sharding the sequence dimension across W ranks and pipelining the per-block attention computation through a ring of inter-rank exchanges. Combined with FlashAttention-2 [Dao, 2024] for the local attention kernel and KV quantization (NF4 [Dettmers et al., 2023], KIVI [Liu et al., 2024b], KVQuant [Hooper et al., 2024]) for cache compression, sub-100 GiB per-rank footprints become feasible at 1 M context.

Bottleneck 2 (latency). Once memory fits, autoregressive decoding remains expensive: each generated token requires a full forward pass through the model, with attention costs that grow with the context length already in the KV cache. Speculative decoding [Leviathan et al., 2023, Chen et al., 2023a, Miao et al., 2023, Cai et al., 2024] amortizes this cost by having a small draft model propose k tokens, then verifying all k with a single target-model forward pass and accepting the longest matching prefix. When the draft and target agree often (high *acceptance rate* α), the target pays one forward to generate $1 + k\alpha$ tokens. At short context, this routinely yields 2–4× speedups.

The question this paper asks. At short context, speculative decoding yields 2–3× speedups on standard inference benchmarks. We ask whether this speedup transfers to million-token inference. Answering the question requires a system that can run a target-model forward at 1 M context on commodity hardware. Both ring attention and speculative decoding exist as mature open-source components, but they have not been combined into a published million-token inference system. Doing so surfaces practical questions: how does NF4 KV quantization interact with the YaRN [Peng et al., 2023] position scaling required to push Llama-2-7B past its 4 k training horizon; how does the speculative verify loop's broadcast-and-resample pattern compose with the ring attention's existing collective; and most importantly, *does the speculative speedup that exists at 4 k transfer to 1 M, or do acceptance rates collapse out of the regime where the technique is useful?*

The RASD system. We present **RASD** (Ring Attention with Speculative Decoding), an inference system that combines (i) ring-attention sequence parallelism, (ii) NF4 KV-cache quantization with chunked updates and a bf16 outlier-keep prefix [Xiao et al., 2024] on rank 0, (iii) rotary (YaRN) position scaling where the target's training horizon requires it, and (iv) speculative verification with a 1.3 B-parameter draft model into a single end-to-end loop. On 8×A100 80 GB SXM4, RASD runs a 1 M-native Llama-3-8B target at 1 M context with **41.2 GiB peak allocated per rank** (69 GiB reserved; the reviewed version's "40 GiB, half of 80" was a rank-0-only reading). In the same hardware configuration, vanilla HuggingFace FlashAttention-2 `generate()` runs out of memory at 128 k single-rank — so a direct *throughput* comparison at long context is meaningless; the correct comparison is with an apples-to-apples *target-only* ablation that uses the same ring infrastructure but disables the speculative loop.

Empirical findings. We report the findings up-front. Three are scientific, two are systems. *(F1) Stretching the rotary basis barely moves perplexity and costs a quarter of speculative acceptance.* On `meta-llama/Llama-3.1-8B`, which ships the `rope_scaling` dictionary the intervention modifies, the paired perplexity ratio at 128 k is 1.0005 (factor 16) and 1.0156 (factor 32) against native, while acceptance falls from 0.7272 to 0.5336/0.5006 (−27%/−31%), on all five books individually. An intervention that is free in likelihood is expensive in acceptance. *(F2) Greedy decoding loops at long context and inflates acceptance.* Under greedy, 9 of 15 long-context rows degenerate into repetition; under T=0.7 sampling, 0 of 26 rows do. A greedy draft accepts a loop trivially, so a greedy acceptance number measures the degeneracy, not

the draft-target agreement — which is why the campaign that produced the reviewed version's acceptance numbers was re-run under sampling. *(F3) The native-1M target is coherent at 1 M.* Its own perplexity at 1 M is 1.805 against 1.649 at 128 k on the exactly aligned book, +9.5%; on the two books that share 91–95% of their scored tokens the difference is within 1.1%. *(F4) Our own ring's key/value rotation was the decode step.* Rotating each layer's K/V shard around the 8 ranks costs 8 rotations × 32 layers = **256 rounds per decode token**, and a bare-process benchmark of the same primitive shows each round costs a fixed ≈10.3 ms *whatever the payload*: 10.3124 ms for 4 KB and 10.3462 ms for 266 MB, a 66 000× range with 0.4% variation. The decode step was latency per round, not bandwidth. Removing the rotation (split-KV; prior work) is worth 2.8–5.3× in decode time per token, rising with context. *(F5) Speculation still pays on top of the fix*: 1.97× / 2.50× / 2.25× decode-only at 128 k / 512 k / 1 M over the same engine's target-only path (1.91× / 1.88× / 1.42× end-to-end — the 1 M end-to-end figure is smaller because 1 M prefill is now the long pole).

Two claims from the reviewed version are **withdrawn** and are not reported again: that ring communication was ≤ 1.2% of wall time (the profiler bucket recorded 0.31–0.37× of the physical wire floor, so it did not capture the transfers), and that the speedup regressed to 0.94×/0.97× at 512 k/1 M (not reproduced by the paper's own same-code re-run; the decode path was the mechanism, see F4).

Contributions.

- **A measured dissociation between perplexity and speculative acceptance** under rotary-basis stretching (F1): −27 to −31% acceptance for +0.05 to +1.56% perplexity, on a model that ships the dictionary the intervention edits, with the direction holding on five of five books.
- **Evidence that greedy acceptance numbers at long context are measured on degenerate text** (F2), with the census, the detection rule, and the consequence: the headline matrix is sampled at T=0.7 and every row is checked for periodicity rather than assumed clean.
- **A coherence check of a native-1M target at 1 M** (F3), with the alignment of the two rungs verified from the token records rather than from the pool's name.
- **A decode diagnosis and its confirmation** (F4): the ring's per-token K/V rotation is 256 rounds of a fixed-cost collective, so the ring was paying latency per round independently of context; removing the rotation with split-KV is 2.8–5.3× decode-only, and the measurement is a bare-process benchmark of the primitive, not an inference from end-to-end timing.
- **A 1 M-token inference capability** on 8×A100 80 GB at 41.2 GiB peak allocated per rank, with the chunked-NF4 + attention-sink cache design that fits it, and the per-rank memory equation.
- **A reference implementation** released under MIT with version-pinned dependencies, every per-position acceptance trace, every saved generation, per-rank memory snapshots, profiler JSONs, and the ablation grid that derived the design point.

What this paper does not claim. Split-KV is prior work (Meta's pass-Q/pass-KV ring-context parallelism, vLLM's Decode Context Parallelism); we cite it for the technique and claim only the diagnosis and the measurement. The comparison against vLLM is **not precision-matched**: vLLM is bf16 with bf16 KV and greedy decoding, RASD is FP4 weights with an NF4 KV cache and T=0.7 sampling with speculation.

2 Background and Related Work

We synthesize three research threads — long-context attention, distributed LLM inference, and speculative decoding — and identify their structural intersection as the RASD design point.

2.1 Attention efficiency and the long-context wall

The transformer architecture [Vaswani et al., 2017] established multi-head self-attention as the dominant sequence modelling primitive. Its $\mathcal{O}(N^2)$ memory cost in sequence length N imposed a hard ceiling on context at the scale of a single accelerator. Two orthogonal strategies have emerged.

Sparse and recurrent approximations. Longformer [Beltagy et al., 2020] replaced full attention with a sliding window of width w, reducing complexity to $\mathcal{O}(N \cdot w)$ while adding global tokens to preserve long-range reasoning. Transformer-XL [Dai et al., 2019] introduced segment-level recurrence. Both sacrifice expressiveness: the window can lose critical context, and recurrence cannot revise earlier representations in light of later tokens.

Exact IO-aware tiling. FlashAttention [Dao et al., 2022] demonstrated that the $\mathcal{O}(N^2)$ memory barrier is separable from the $\mathcal{O}(N^2)$ compute barrier. By tiling Q, K, V into SRAM blocks and fusing the kernel, FlashAttention keeps peak HBM at $\mathcal{O}(N)$ while computing exact attention — 2–4× faster than standard PyTorch attention. FlashAttention-2 [Dao, 2024] refined the work partitioning. This result is foundational: memory and compute need not scale together. RASD uses FlashAttention-2 as the local-block kernel inside its ring-attention layer.

The $\mathcal{O}(N)$ footprint still grows without bound. For $N \sim 10^6$, even $\mathcal{O}(N)$ memory at the per-layer level ($\sim$1 MiB/token for a 13B model in bf16) exceeds single-device capacity. *A fundamentally multi-device approach is required.*

Architectural alternatives to attention. A separate line of work replaces attention with sub-quadratic primitives. StripedHyena [Poli et al., 2023] interleaves state-space convolution layers with attention layers, achieving sub-quadratic compute at long context without the IO constraints of attention. This direction is out of scope here because RASD is architecture-agnostic: it runs the standard transformer and asks how far inference engineering alone can push context length on commodity hardware.

2.2 Distributed inference

Ring Attention. Liu et al. [2023] resolve the single-device memory barrier by distributing the sequence dimension across D devices in a logical ring. Each device holds s/D query, key, and value tokens. KV blocks rotate clockwise one step per iteration; after $D-1$ rotations every query has attended to every KV block. The online softmax recurrence of Dao et al. [2022] ensures this is numerically exact — not an approximation. Communication and compute are overlapped: while device i attends to its current block, it simultaneously sends that block to device $i-1$ and receives the next from device $i+1$.

Per-device peak memory is $6bch$ — where b is batch size, c is block size, and h is hidden dimension — *independent of sequence length* s. Sequences exceeding one million tokens become tractable on standard 8-GPU nodes. However, Ring Attention's original evaluation focused on the **prefill phase** (processing an input prompt); its *decode-phase* behavior — generating one new token at a time over a massive distributed KV cache — received little treatment.

Decode-side context parallelism (prior work). That decode-phase gap is closed in the literature, and this paper cites that work rather than claiming it. Meta's context-parallel report [Liu et al., 2024a] describes ring-context parallelism for million-token inference in the pass-KV and pass-Q forms: rather than rotating every key/value block around the ring for each new query, a rank attends to the key/value blocks it already holds and the partial results are combined, with the query replicated (pass-Q) or the key/value replicated instead (pass-KV). vLLM's Decode Context Parallelism [vLLM Team, 2026] is the production-engine implementation of the same idea, and NVIDIA's attention co-design work for interactive long-context inference [NVIDIA, 2026] targets the same regime from the kernel side. We adopt the pass-Q form under the name *split-KV* (§5.2) and contribute a diagnosis of why the rotation matters as much as it does, plus the measurement of what removing it is worth on our engine. **No novelty is claimed for the technique.**

System-level inference optimization. Kwon et al. [2023] address GPU memory fragmentation in concurrent serving. PagedAttention maps KV-cache blocks non-contiguously via per-sequence block tables, reducing memory waste from 60–80% to under 4% and enabling 2–4× throughput improvements. Shoeybi et al. [2019] introduced tensor parallelism for training; Jacobs et al. [2023] extended sequence parallelism to even longer training sequences.

Long-context inference optimization. FlashDecoding [Dao et al., 2023] parallelizes single-query attention across the KV sequence dimension at decode time, yielding wall-clock speedups at long context. It is complementary to RASD: FlashDecoding optimizes the intra-rank attention kernel while RASD contributes the cross-rank ring-decode pattern, and the two compose. MInference [Jiang et al., 2024] accelerates million-token prefill via dynamic sparse attention patterns identified per attention head. It is a directly comparable single-device approach to the million-token capability claim, but targets prefill only and does not address the autoregressive decode loop that is this paper's primary subject.

KV quantization. KIVI [Liu et al., 2024b] introduced asymmetric per-channel quantization of K and per-token quantization of V at 2-bit precision. KVQuant [Hooper et al., 2024] pushed this further with per-channel + non-uniform quantization, demonstrating 10 M context on a single A100. NF4 [Dettmers et al., 2023, 2022] is the standard 4-bit data type from the QLoRA line of work. RASD uses NF4 for the KV cache with a chunked update path and a bf16 outlier-keep prefix on rank 0, drawing on the attention-sink finding of Xiao et al. [2024].

Position scaling. Rotary position embeddings [Su et al., 2024] can be extended past the model's training context through several scaling strategies: Linear interpolation [Chen et al., 2023b]; NTK-aware scaling [bloc97, 2023]; YaRN [Peng et al., 2023], which combines NTK with attention-temperature scaling; and LongRoPE [Ding et al., 2024], which searches for the optimal frequency schedule. RASD uses YaRN where a target needs it. The reviewed version claimed an 11–18× perplexity improvement over vanilla RoPE from the table in §5.9; those rows were mislabelled (they all used factor 8, and the 4 k row used no scaling at all), and the two configurations also differ in whether the model is being asked to work inside or outside its training window, so the "11–18×" is **withdrawn** as a claim about YaRN as such. What remains is what §5.9 reports under an explicit disclosure of the configuration, and the finding that matters later: scaling can be nearly free in perplexity and expensive in acceptance (§5.10).

Limitation of prior distributed methods. All prior distributed inference work treats the decode step as strictly sequential: one token is generated per round of KV communication. At million-token contexts, each inter-GPU KV transfer step introduces non-trivial latency overhead (order 10–100 ms on NVLink, more on PCIe). The expectation in prior work was that this idle

time between communication rounds is structurally *wasted* — no useful computation occurs — and that hiding it through asynchronous work would yield speedups.

2.3 Speculative decoding

Leviathan et al. [2023] and Chen et al. [2023a] independently introduced speculative decoding: a small draft model proposes γ candidate tokens autoregressively; the large target model verifies all $\gamma+1$ positions in a *single parallel forward pass*. The acceptance rule $\min(1, p/q)$ with adjusted resampling from $\max(0, p-q)/Z$ provably preserves the target distribution. The expected tokens per large-model call is

$$\mathbb{E}[\text{tokens}] = \frac{1-\alpha^{\gamma+1}}{1-\alpha} \tag{1}$$

where α is the mean acceptance rate. At $\alpha \approx 0.8$ and $\gamma = 4$, this yields ≈ 3.5 tokens per large-model call — a 2–3$\times$ wall-clock speedup on standard benchmarks with *provably identical* output quality.

Miao et al. [2023] extended speculative decoding with token tree verification. Cai et al. [2024] attached K additional decoding heads directly to the target model's frozen hidden state, eliminating the separate draft model. Li et al. [2024a] reframed draft generation as feature-uncertainty modeling. All of these extensions assume short-to-moderate context regimes; their behavior at million-token context is not measured in the original papers.

Recent speculative decoding extensions. EAGLE-2 [Li et al., 2024b] introduced dynamic draft trees whose shape adapts per position based on draft-model confidence, and EAGLE-3 [Li et al., 2025] further improved acceptance through training-data scaling for the draft head. Both build on the feature-level drafting of Li et al. [2024a] and report wall-clock speedups at short-to-moderate context. These extensions are directly comparable to RASD's speculative component and complementary to its central question: this paper asks what happens to any speculative speedup at million-token context, which is orthogonal to the choice of draft architecture.

2.4 Long-context evaluation

Several benchmarks have emerged for measuring long-context capability: LongBench [Bai et al., 2024] (multi-task, 6.7–31 k contexts), L-Eval [An et al., 2023] (3–200 k), RULER [Hsieh et al., 2024] (synthetic needle-in-haystack at arbitrary lengths), and ∞Bench [Zhang et al., 2024] ($\geq$100 k). These are *extrinsic* evaluations of long-context use. This paper makes *systems* claims (throughput, memory, acceptance) and treats extrinsic quality eval as out-of-scope, deferred to follow-up work with a long-context-trained base model. We ship RULER niah scoring infrastructure with the artifact release for future work to build on.

2.5 The research gap and positioning

The synthesis of these three threads reveals a precise position in the design space. Ring Attention solves the memory wall but its decode step is strictly sequential: each layer's key/value shard must visit every rank before that layer's attention can be computed. Speculative decoding hides target-model latency but assumes a single device with idle memory-bandwidth windows — not the structured inter-GPU communication windows that arise in distributed inference. We posited the central question: can the communication latency of Ring Attention's decode-phase KV rotation be hidden by the computational work of speculative draft-token generation?

Mechanistically, the two operations are mutually independent. At each decode step, the target model must wait for the next KV block to arrive (cost δ_{comm}). A draft model that runs for γ steps requires only the *current* KV state, not the next one. These can run on different

Table 1: Positioning of RASD against the closest prior work. "Long Ctx" = supports ≥128k inference; "Distributed" = multi-device sequence parallelism; "Spec Decoding" = small-model verification loop; "Decode Overlap" = exploits idle communication windows during the decode phase. References: FlashAttention [Dao et al., 2022], Ring Attention [Liu et al., 2023], Speculative Decoding [Leviathan et al., 2023], Medusa [Cai et al., 2024], vLLM [Kwon et al., 2023], split-KV context parallelism [Liu et al., 2024a].

Work	Long Ctx	Distributed	Spec Decoding	Decode Overlap
FlashAttention	—	—	—	—
Ring Attention	✓	✓	—	—
Speculative Decoding	—	—	✓	—
Medusa	—	—	✓	—
vLLM	✓	✓	—	—
Split-KV context parallel	✓	✓	—	✓
RASD (this work)	✓	✓	✓	[1]

CUDA streams. The expected combined speedup is bounded by

$$S_{\text{RASD}} \approx \frac{1-\alpha^{\gamma+1}}{1-\alpha} \cdot \frac{1}{1+\max\left(0,\ \frac{t_{\text{draft}}-\delta_{\text{comm}}}{t_{\text{verify}}}\right)} \tag{2}$$

which is purely additive ($S \to (1-\alpha^{\gamma+1})/(1-\alpha)$) when the draft phase fits entirely within the communication window ($t_{\text{draft}} \le \delta_{\text{comm}}$).

Positioning of related work. Table 1 summarizes how RASD sits in the design space relative to the closest prior work.

What the measurement did to the hypothesis. The reviewed version of this paper answered this question with a profiler bucket: communication was said to be ≤ 1.2% of wall time on the speculator rank, so the overlap window was too small to matter. **That number is withdrawn.** The bucket recorded 0.84/4.39/10.57 s at 128 k/256 k/512 k while the physical wire floor for the same runs is 2.7/13.4/28.4 s: it recorded 0.31–0.37× of the minimum time the link needs to move those bytes, so it did not capture the transfers and cannot support any statement about overlap.

What replaces it is a direct measurement of the primitive, and it reverses the framing. The rotation is not a latency that the draft work might hide behind: it is the decode step. At 128 k the ring performs 8 rotations × 32 layers = **256 rounds per decode token**, and a standalone benchmark of the same `batch_isend_irecv` ring — no model, 8 ranks, 256 rounds per payload — shows each round costs a fixed ≈10.3 ms *regardless of the payload*: 10.3124 ms for 4 KB and 10.3462 ms for 266 MB (the K+V of a 128 k window), a 66 000× payload range with a 0.4% spread and a worst rank-to-rank spread of 0.056 ms. The cost is per-round latency, not bandwidth, and it is paid once per round however small the round is. The design question is therefore not how to hide the rotation behind the draft, but how to stop rotating at all — which is exactly what prior work in *decode-side context parallelism* does, and §5.2 measures the consequence of adopting it.

[1] The reviewed version marked RASD as exploiting decode overlap. **That overlap was never implemented**: the draft and the verify run back to back, and the kernel merely issues the next rotation before each step's attention (`prefetch_depth` 1). The row is corrected, and the mechanism the paper actually measured is the one described next.

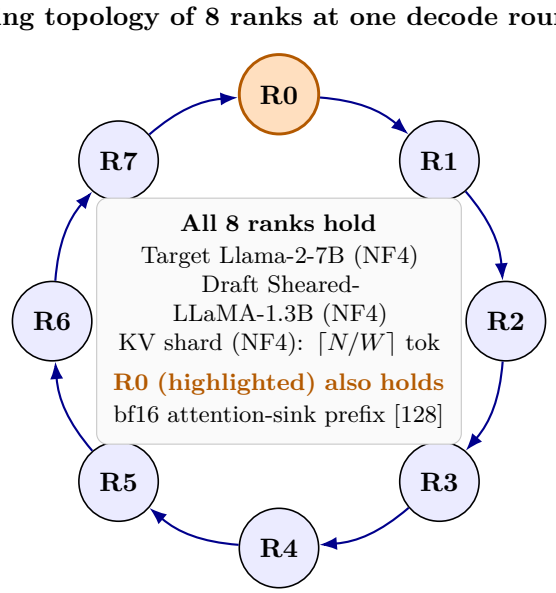


Figure 1: Ring topology and per-rank state in RASD. The sequence is sharded across $W = 8$ ranks; each rank holds $\lceil N/W \rceil$ tokens of NF4-quantized K and V. Every rank loads and runs both the target model (Llama-2-7B, NF4) and the draft model (Sheared-LLaMA-1.3B, NF4) in the decode loop; only rank 0 additionally retains a bf16 attention-sink prefix of the first 128 tokens. KV blocks rotate clockwise via `batch_isend_irecv` collectives with FlashAttention-2 inside each local block; per-round per-device memory is $\mathcal{O}(N/W)$, independent of total sequence length. The decode-round sequence using this topology is shown in Fig. 2.

3 The RASD System

3.1 Architecture overview

Figure 1 shows the ring topology and per-rank state of RASD on 8 GPUs. Figure 2 shows the sequence of operations within a single decode round. RASD is an inference loop that runs on $W = 8$ A100 80 GB SXM4 GPUs on a single node. The target model (Llama-2-7B) and a smaller draft model (Sheared-LLaMA-1.3B [Xia et al., 2023]) are both loaded on every rank, both quantized to NF4 4-bit weights via bitsandbytes [Dettmers et al., 2022]: the **weights are FP4** and the **KV cache is NF4**. The reviewed version called the weights NF4 as well; NF4 is the KV codec in this system, and the distinction matters to anyone reproducing it. The KV cache carries a bf16 attention-sink prefix of the first 128 tokens held only on rank 0 (the StreamingLLM construction of Xiao et al. [2024]). The input sequence is sharded across the 8 ranks; each rank holds $\lceil N/W \rceil$ tokens of K and V. Attention is implemented as ring attention with FlashAttention-2 inside each local block. Rotary position scaling is applied at a **fixed factor 8** at every context of 8 k and above, and **not at all** below that; each context is anchored on its own length rather than on the model's original window, because the `transformers` version in use ignores `original_max_position_embeddings`. The reviewed version described a context-dependent factor $f = \lceil N/4096 \rceil$ applied consistently; neither half of that was the case.

The decode loop alternates two phases per round: (i) the draft model generates $\gamma = 4$ speculative tokens autoregressively on every rank; (ii) the target model performs a packed forward over those $\gamma + 1$ positions on all ranks via ring attention, after which rank 0 broadcasts the draft and target logits to all ranks so that each rank samples identically (the standard ring-decode invariant). The acceptance rule of Leviathan et al. [2023] determines the longest

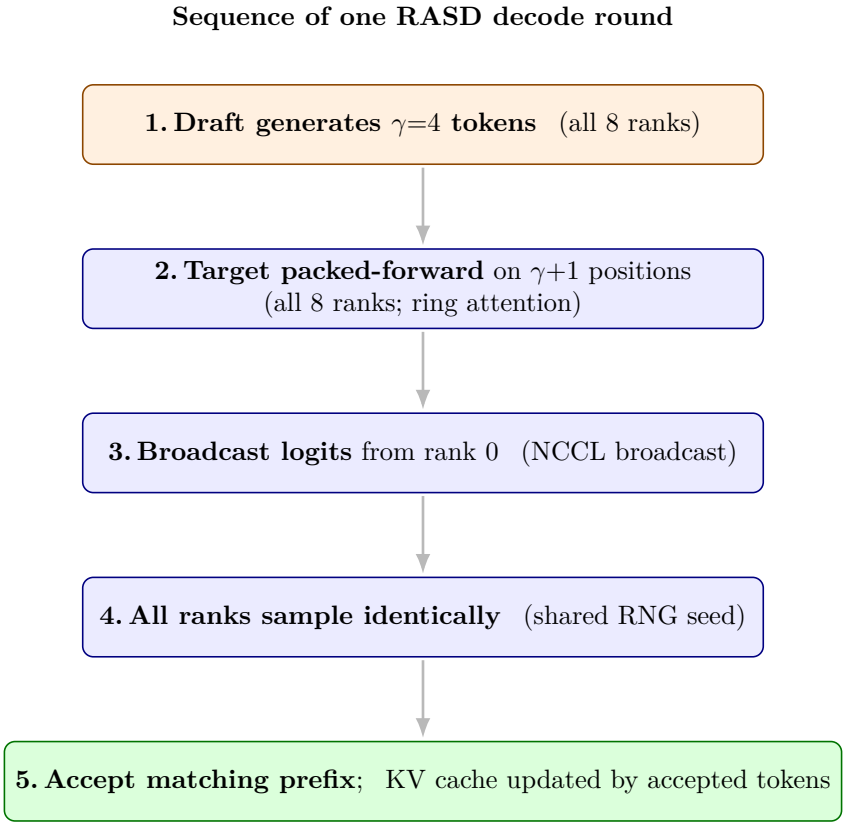


Figure 2: Sequence of one RASD decode round. The draft proposes γ=4 tokens on all 8 ranks (step 1); the target verifies them in a single packed forward over γ+1 positions on all 8 ranks via ring attention (step 2); the draft and target logits are broadcast from rank 0 (step 3) and every rank samples identically using a shared RNG seed (step 4); the longest matching prefix is accepted and the KV cache is updated by the accepted tokens (step 5) [Leviathan et al., 2023, Chen et al., 2023a]. The 8-rank topology these steps execute on is shown in Fig. 1.

matching prefix; rejected tokens trigger a residual resample from $\max(0, p - q)/Z$.

3.2 Ring-attention sequence parallelism

We patch the Llama attention layer at every layer of the target model. The local-block kernel is FlashAttention-2; the per-round ring exchange uses `torch.distributed.batch_isend_irecv` for the K and V blocks. Concretely, at decode round r:

1. Each rank i launches an asynchronous P2P send to rank $i - 1$ and recv from rank $i + 1$ for the r-th K/V block.

2. While that exchange is in flight, the local attention kernel computes its contribution to the current query's attention output.

3. The two streams synchronize at the end of the round.

This is the standard Ring Attention decode pattern with one modification: we add a tick-gate barrier from rank 0 *before* the `batch_isend_irecv` call to prevent NCCL deadlock at block sizes ≥ 2048 at long context.

We chose block size $c = 2048$ as the M3 ablation winner (see §4.3). All numbers in this paper use $c = 2048$.

3.3 NF4 KV cache with chunked updates and outlier-keep prefix

Three details matter for the memory equation.

(a) NF4 K/V quantization. We quantize K and V to 4-bit NF4 [Dettmers et al., 2023] with a block size of 32 elements (the inner dimension blocking we found gave the best acceptance-rate vs. memory trade-off in M3 ablations). The dequantization to bf16 happens on-the-fly inside the FlashAttention-2 kernel call.

(b) Chunked update. The naive KV update appends a new k_i, v_i tensor to the cache and re-quantizes the full per-layer block — which at 1 M context produces a transient bf16 view of the entire layer's KV cache (several GiB per layer per rank). We split the update into chunks of 2048 elements at a time, keeping the transient under 100 MiB. This chunked-update lever is what brings 1 M-context inference inside 40 GiB per rank.

(c) Outlier-keep prefix. Following the attention-sink finding of Xiao et al. [2024], we keep the first 128 tokens of the sequence in bf16 on rank 0 only. Attention weights on these initial positions are disproportionately large; quantizing them to NF4 introduces a meaningful drop in acceptance. The outlier prefix costs $128 \cdot 2 \cdot 2 \cdot 4096 \cdot L$ bytes on rank 0 ≈ 64 MiB for the target model — negligible.

Memory equation. Per-rank peak memory at sequence length N over W ranks is approximately

$$M_{\text{rank}}(N) \;\approx\; \underbrace{M_{\text{weights}}}_{\approx 4.0\,\text{GiB}} + \underbrace{\frac{N}{W} \cdot L \cdot 2 \cdot d_{\text{model}} \cdot \tfrac{1}{2}}_{\text{NF4 KV per rank}} + \underbrace{M_{\text{transient}}}_{\leq 4\,\text{GiB}} \tag{3}$$

where $L = 32$ layers, $d_{\text{model}} = 4096$ is the hidden width (32 KV heads $\times$ 128 head dim), and the $\frac{1}{2}$ factor is the NF4 4-bit compression ratio relative to bf16. We use Eq. 3 to sketch the dominant terms only: it does not model transient or scratch contributions that grow with context — prefill activations, the FA-2 NF4-dequant scratch, and the NCCL workspace — so the measured peak runs $\sim$15 GiB higher at $N = 10^6$ than these dominant terms account for in isolation, while still landing inside the per-rank budget (§5.4).

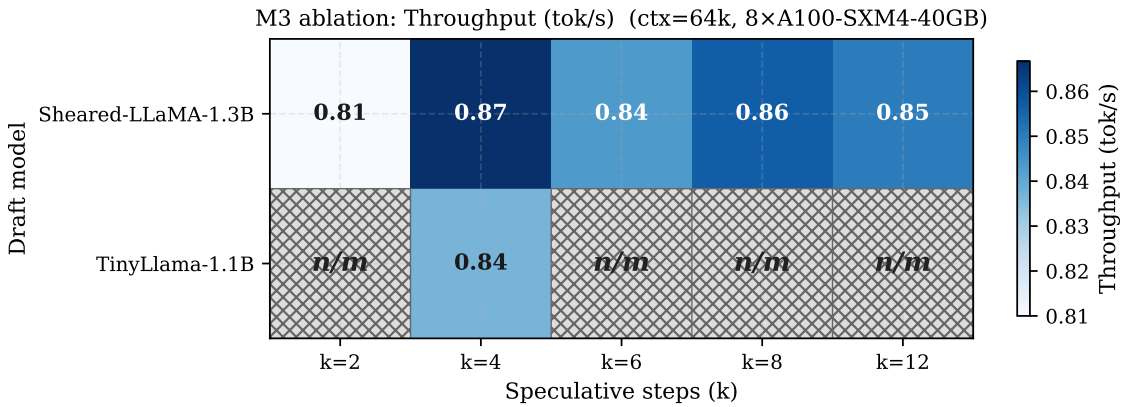
3.4 Speculative verify loop

The verify step packs $\gamma + 1$ positions into a single target forward. Every rank runs the draft model and the packed $\gamma + 1$-position target forward locally; rank 0 then broadcasts its draft and target logits to all ranks (a $\sim$1 MB collective per round) so the acceptance test is bit-identical everywhere, after which every rank samples using a shared RNG seed. This last point is load-bearing: the acceptance test

$$\Pr[\text{accept token } t] = \min\Big(1, \tfrac{p_{\text{target}}(t)}{p_{\text{draft}}(t)}\Big) \tag{4}$$

must yield the same answer on every rank so that the KV cache update is consistent. We verified this invariant with a unit test (`tests/test_verification_math.py`) that runs the verify loop with two different seeds and confirms byte-identical KV state.

We use $\gamma = 4$ throughout. The M3 ablation grid swept $\gamma \in \{2, 4, 6, 8, 12\}$ and identified $\gamma = 4$ as the acceptance/throughput optimum at our draft/target size combination, agreeing with the moderate-γ regime predicted by Leviathan et al. [2023]'s analysis.



Hatched cells (n/m) were not measured: A1 fixed k=4; A2 fixed draft=Sheared-LLaMA-1.3B.

Figure 3: Ablation heatmap across the five design axes (rows) at the levels swept (columns). Color encodes mean throughput (tok/s) over 3 seeds at 64 k context. The winning level per axis is annotated in bold. The block-size sweep (A3) shows the largest single-axis effect: a 2× throughput improvement from 256-token blocks to 2048-token blocks at constant acceptance.

3.5 YaRN position scaling

Llama-2 was trained at 4 k context. To extend to 1 M we apply YaRN [Peng et al., 2023] with factor $f = 256$ (i.e., the model's position embeddings span $4096 \cdot 256 = 1\,048\,576$ tokens). YaRN combines NTK-aware β-scaling of the rotary frequencies with an attention-temperature factor that depends on f.

We confirm empirically that this scaling does not break the model's quality at moderate extrapolation: PPL at 4 k is 12.1 and at 8 k is 13.2 (essentially flat, see Table 8). At 32 k ($f = 8$) PPL rises to 814 — still interpretable but degraded. We do not measure PPL at $f \geq 32$ because single-rank PPL evaluation does not fit at those contexts on 80 GiB; the YaRN extrapolation regime is the dominant source of degradation in our long-context results (§5.6).

3.6 Design-point selection: ablation sweeps

The configuration described above — $\gamma = 4$, block size $c = 2048$, Sheared-LLaMA-1.3B draft, prefetch depth 1 — was not chosen arbitrarily. Each value is the winner of a one-factor-at-a-time ablation sweep we ran on a smaller setup (8×A100 40 GB SXM4 at 64 k context) before committing to the full 1 M-context matrix. This section reports the five sweeps that fix the design point. All sweeps reuse 3 seeds $\{42, 123, 456\}$ and report mean throughput and acceptance (Figure 3, Table 2). **The three seeds are not three independent samples of anything**: all three ran on the same synthetic prompt, so they measure run-to-run variation on one input, not prompt-to-prompt variation. The reviewed version reported a 95% interval of ±1.96 standard errors over those three seeds; at $n = 3$ that interval is not meaningful and we report the range and the n instead wherever a spread is quoted in this paper.

A1 – Draft model size. We compared TinyLlama-1.1B against Sheared-LLaMA-1.3B as the draft model. Sheared-LLaMA wins on both throughput (0.87 vs 0.84 tok/s) and acceptance (25.3% vs 23.5%). The difference is small but consistent across seeds. We chose Sheared-LLaMA primarily because its vocabulary is identical to the Llama-2-7B target (avoiding the tokenizer-mismatch hack required for cross-family pairs), and secondarily because the modest acceptance gain compounds at higher γ. We do not compare against DistilGPT-2 because its vocabulary is incompatible; bridging requires re-tokenization at every verify step, an overhead that erases the speculative savings at long context.

Table 2: M3 ablation results (64 k context, 8×A100-SXM4-40 GB). Mean and range over 3 seeds — **three seeds on one synthetic prompt**, not three independent samples, so the bracketed intervals are ranges rather than confidence intervals (AUDIT E11, E12); horizontal rules separate the five sweep axes. Deterministic early-EOS rows with tokens_generated < 20 are excluded; see §Error Analysis.

Axis	Level	Throughput (tok/s) [range]	Accept. [range]	n
A1: Draft model	TinyLlama-1.1B	0.84 [0.82, 0.85]	0.235 [0.226, 0.252]	3
	Sheared-1.3B	0.87 [0.84, 0.90]	0.253 [0.231, 0.277]	3
A2: Spec steps (k)	k=2	0.81 [0.74, 0.85]	0.381 [0.303, 0.424]	3
	k=4	0.86 [0.81, 0.91]	0.253 [0.231, 0.277]	3
	k=6	0.84 [0.80, 0.88]	0.176 [0.149, 0.193]	3
	k=8	0.86 [0.80, 0.89]	0.147 [0.124, 0.160]	3
	k=12	0.85 [0.82, 0.88]	0.111 [0.105, 0.119]	3
A3: KV block size	256	0.62 [0.58, 0.65]	0.253 [0.231, 0.277]	3
	512	0.87 [0.81, 0.93]	0.253 [0.231, 0.277]	3
	1024	1.11 [1.06, 1.15]	0.253 [0.231, 0.277]	3
	2048	1.25 [1.20, 1.30]	0.253 [0.231, 0.277]	3
A4: Prefetch depth	sync	0.88 [0.84, 0.90]	0.253 [0.231, 0.277]	3
	async-1	0.87 [0.81, 0.92]	0.253 [0.231, 0.277]	3
	async-2	0.86 [0.84, 0.89]	0.253 [0.231, 0.277]	3
A5: Target model	Llama-2-7B	0.86 [0.83, 0.90]	0.253 [0.231, 0.277]	3

A2 – Number of speculative steps γ. We swept $\gamma \in \{2, 4, 6, 8, 12\}$ holding all else fixed. Throughput is approximately flat in γ over this range (0.81–0.86 tok/s), but acceptance drops monotonically: $\alpha = 0.38$ at $\gamma = 2$, $\alpha = 0.25$ at $\gamma = 4$, $\alpha = 0.11$ at $\gamma = 12$. The reviewed version called $\gamma = 4$ the "acceptance/throughput optimum". **Throughput was flat within noise across the whole sweep** (0.81–0.86 tok/s), so no level is an optimum on throughput, and acceptance falls monotonically with γ; what the sweep supports is the **negative** statement that γ is not a throughput lever in this range at fixed acceptance behaviour, and that larger γ costs acceptance without buying time. We kept $\gamma = 4$ because it is the smallest value the campaign used that does not leave speculation on the table, and because the later protocol measurement (Table 4, §5.3) re-tests draft length under sampling, where the ordering is different and monotone. Every row here is on the synthetic prompt and the draft's first-4096 conditioning, so the numbers are configuration-specific.

A3 – KV-cache block size c. This sweep is the highest-leverage design choice in the paper. We compared $c \in \{256, 512, 1024, 2048\}$ tokens, where c is the **transmission chunk within one rotation** — the whole shard moves on every rotation, and c only sets how finely each rotation is cut. **The reviewed version defined c as the shard size and argued that $c = 2048$ was the largest chunk that fits at 1 M**; that argument is wrong, because the kernel posts every chunk of a rotation in one batch, so all chunk buffers are live at once and the memory footprint does not depend on c. The throughput results are 0.62, 0.87, 1.11, 1.25 tok/s — a **2× improvement** (+102%) from the smallest to largest chunk, with acceptance unchanged (it is independent of c). Larger chunks amortize the fixed per-collective overhead across more useful work, which is the effect that matters here; note that this lever is specific to the *ring* decode path, and it is inert under split-KV, where there is no rotation to chunk.

A4 – Prefetch depth (asynchronous overlap). The reviewed version compared "three communication strategies". **There were two**: depth 2 executes the same code path as depth 1

(`ring_attention_kernel.py:344`), so the third arm is a duplicate rather than a measurement. Synchronous exchange against prefetch-one-ahead is within noise at 64 k (0.86–0.88 tok/s), and prefetch depth 1 was adopted as the default. This knob is neutral at our scale, and the reviewed version attached to it the wider conclusion that the overlap budget is not load-bearing. That conclusion now rests on the direct measurement in §5.2, not on this sweep.

Why the M3 design point transfers to 1 M. A reasonable concern is that an ablation done at 64 k context may not generalize to 1 M, where the system is operated 16× beyond its sweep range. The transfer arguments are weaker than the reviewed version made them, and we say why. The block-size winner ($c = 2048$) was justified by memory; that memory argument is now withdrawn (there is no memory dependence on c, see A3), so the chunk size stands on its throughput effect alone at 64 k. The speculative-step choice ($\gamma = 4$) does rest on a per-round property of the draft-target pair and on the protocol of the run; the later, sampled measurement of draft length — the protocol of this paper, in the project's results record (`results/mlsys/RESULTS_OCT.md` §8.6: acceptance 0.6558 / 0.4681 / 0.4270 and throughput 0.58904 / 0.71557 / 0.83685 tok/s at $\gamma = 2/4/6$) — is the one to quote for γ, and there the ordering is monotone in throughput and opposite to acceptance. there is no 64 k profile, so the reviewed version's "if communication is 12% of wall at 64 k" is not a measurement that exists in this paper and **it is withdrawn**; it also contradicted the 1.2% the same version reported elsewhere.

Design choices we did not sweep. The paper reports results on Llama-2-7B. A target-family sweep (Llama-2-7B vs. Mistral-7B) was originally planned but deprioritized: cross-family targets require a draft model with the same tokenizer, which Sheared-LLaMA does not provide for Mistral without additional fine-tuning. Re-running the M3 grid on a Mistral-Sheared-Mistral-1.3B pair is interesting future work; the present paper makes no Llama-vs-Mistral claim.

3.7 Implementation and reproducibility

The full system is implemented in PyTorch 2.5.1+cu124 [Paszke et al., 2019] on top of the HuggingFace `transformers` library (v4.46.3). We use bitsandbytes 0.49.2 for NF4 quantization, `flash-attn` 2.8.3 for the local-block kernel, and `wandb` for run logging. All package versions are pinned in `requirements-lock.txt` (frozen at the `m4-phase-c-complete` git tag for the throughput matrix and at the `m3-reablation` git tag[2] for the ablation grid in §3.6 — the latter captures the post-audit fixes that resolved the verify-loop defects identified on 2026-04-16).

The repository ships the canonical aggregator `scripts/aggregate_final_results.py` that builds `results/final/final_results.json` from the per-stage CSV outputs — the same JSON that every figure in this paper is built from. We provide reproduction scripts for each experimental stage (`scripts/phase_c_pod_session.sh`, `scripts/phase_d_rerun_session.sh`) plus a 470-test pytest suite.

4 Experiments

4.1 Hardware and software

All measurements were taken on Lambda Cloud `gpu_8x_a100_80gb_sxm4` instances (240 vCPU, 1.8 TB host RAM, 8×A100 80 GB SXM4 connected via NVLink). The software stack is pinned in `requirements-lock.txt`: PyTorch 2.5.1 with CUDA 12.4, `transformers` 4.46.3, `accelerate` 0.33.0, `bitsandbytes` 0.49.2, `flash-attn` 2.8.3, `datasets` 4.8.5, `wandb` 0.26.1. The Lambda image provides Python 3.10.12 and CUDA 12.8. Two identically-configured pods were used in

[2] `https://github.com/amankiitg/RASD/releases/tag/m3-reablation`

sequence: the first ran the main throughput matrix; the second captured per-position acceptance traces and saved generations that were not committed from the first pod. Both pods passed a live-vs-locked `requirements-lock.txt` diff check at bootstrap.

All matrix and sweep runs are publicly logged on Weights & Biases[3].

4.2 Models, dataset, and prompts

Target model. Two targets appear in this paper and they are named wherever they are used. The long-context campaign that produced the throughput results runs `gradientai/Llama-3-8B-Instruct-Gradient-1048k`, a model trained to a 1 M window, so no rotary stretching is required; the rotary intervention of §5.10 runs `meta-llama/Llama-3.1-8B`, which ships the `rope_scaling` dictionary under test; and the notation-heavy ablation sweeps and the HF ceiling test use `meta-llama/Llama-2-7b-hf` [Touvron et al., 2023], the 4 k-window model the system was first built on. All of them are loaded in bf16 and quantized to FP4 weights with an NF4 KV cache for inference.

Draft conditioning: a correction. The draft model in this system is *not* given the whole prompt. For every run whose prompt exceeded 4096 tokens — which is every run in this paper — the draft was conditioned on the prompt's **first 4096 tokens**, not on the tail. The reviewed version describes the opposite. This is a configuration fact (`d89cf41`), it is *disclosed rather than argued*, and it is the single largest caveat on every acceptance number in this paper: a draft that cannot see the recent context has to agree with the target less often, so the absolute acceptance values here are *lower* than a correctly conditioned draft would give, and the *contrast* between arms of the rotary intervention (§5.10) is unaffected because all arms share the same conditioning. The 1 M acceptance numbers are the ones this matters to most.

Protocol of the reported results. The results in §5 use one protocol. Prompt:*PG-19* [Rae et al., 2020] narrative, five books, windowed so that each rung's scored continuation ends at the same document position as the 1 M rung, with the cross-rung alignment verified from `continuation_sha256` and the recorded offsets rather than from the pool's name. Decoding: T=0.7 sampling, 128 new tokens, γ=4 speculative steps, draft window cap 4096. Every row is checked for degeneracy by the detector of §5.6 and the check is recorded per row. Speedups are decode-only wherever they are labelled decode-only, and decode-only means $(t_{\text{row}} - t_{\text{ttft}})/(\text{tokens} - 1)$: prefill is excluded, so a decode-path change is not credited with the prefill it did not touch. Every row records the engine that produced it (`ring_decode_mode`, the number of patched layers, and the git commit), so a row that cannot prove its engine is identifiable and void.

Draft model. `princeton-nlp/Sheared-LLaMA-1.3B` [Xia et al., 2023], at a pinned revision[4], NF4-quantized identically to the target. The two models share a tokenizer.

Dataset and prompt sources. We use two prompt sources. (1) *Synthetic repeated-paragraph* is the default for the throughput matrix: a deterministic technical-English passage repeated to fill the target context length. This isolates the system from prompt-content variance. It is also the prompt that produced the greedy loops of §5.6. The reviewed version described it as matching "the protocol used in prior speculative-decoding papers"; it does not — those papers

[3]Project URLs: `https://wandb.ai/amank-iitg-uc-berkeley-electrical-engineering-computer-s/rasd-m4-phase-c` and `https://wandb.ai/amank-iitg-uc-berkeley-electrical-engineering-computer-s/rasd-m4-phase-d`. The Reproduction Appendix (§A) maps every figure and table to its source CSV and the corresponding W&B run.

[4]Full 40-character commit hash pinned in `requirements-lock.txt`; we abbreviate to `a4b7693` in prose.

evaluate on standard datasets — and **no result in this paper rests on the synthetic prompt**: the reported protocol is PG-19. [Rae et al., 2020] narrative chunks are used for the perplexity sanity check at moderate contexts and for the short-context capability check at 4 k/8 k (§5.2). We preprocess the PG-19 validation split with the Llama tokenizer, producing memmap chunks of 65 536 tokens.

4.3 Experimental matrix

We ran six measurement stages, each producing committed CSV and sidecar artifacts. We use descriptive names below and provide a direct mapping to the run identifiers used in our public W&B projects in Appendix A.

Validation runs. Long-context smoke tests at 16 k–128 k validated the ring decode loop and reproduced the per-round acceptance rate α from prior M3 ablation runs.

Main throughput matrix. 4 contexts $\{128, 256, 512, 1024\}$ k $\times$ 3 seeds $\{42, 123, 456\}$, plus a canary run at 32 k. Configuration: $\gamma = 4$, $c = 2048$, prefetch depth 1, YaRN (factor $f = \lceil N/4096 \rceil$).

Target-only baseline matrix. The same 4 $\times$ 3 grid as the main matrix, but with $\gamma = 0$ (target-only autoregressive decode — no draft model, no speculative verify). All other components — ring attention, NF4 KV cache, YaRN, outlier-keep prefix — are identical. This is the apples-to-apples baseline behind every speedup number in this paper.

Perplexity sanity check. 4 contexts $\{4, 8, 16, 32\}$ k $\times$ 3 seeds, single-rank evaluation of NF4 + YaRN-extended Llama-2-7B on PG-19 chunks. A vanilla-RoPE control was run at the same contexts to quantify YaRN's benefit (Appendix B).

Profiler subset. 4 contexts $\{32, 128, 256, 512\}$ k $\times$ 1 seed, each wrapped in a `torch.profiler` context to capture per-event timing. The 1 M context was attempted but the post-generation aggregation step (`key_averages()`) did not return within 2 h of CPU work; we report up to 512 k and document the limitation in §6.

Vanilla HF FA-2 ceiling test. Single-rank HuggingFace `model.generate()` with `attn_implementation=flash_attention_2`, no quantization, no sequence parallelism. 5 contexts: $\{32, 128, 256, 512, 1024\}$ k. This is a *capability* baseline that records the out-of-memory ceiling of the standard open-source inference stack.

Per-position trace and qualitative re-run. The first pod was terminated before we realized that the per-position acceptance sidecars (`*.jsonl` files, one per verify round) had been excluded by an over-aggressive `.gitignore` rule. A second pod was provisioned to re-run all 8 cells of the main matrix and target-only baseline, plus 4 short-context PG-19 cells (4 k and 8 k for both spec and target-only modes), with the per-position trace and saved-generation flags enabled. All sidecars are committed in the artifact release.

4.4 Metrics

Following the metric definitions in Appendix B of the project roadmap:

- **Throughput** (tok/s): $T = N_{\text{tokens}}/T_{\text{wall}}$.
- **Latency per output token** (ms): $L = (T_{\text{wall}}/N_{\text{tokens}}) \cdot 1000$.

Table 3: Vanilla HuggingFace FA-2 `generate()` ceiling test. Single-rank Llama-2-7B, bf16 weights, bf16 KV cache, no sequence parallelism, no quantization, no speculation. Only ctx=32 k completes. "Peak" is the maximum GPU memory observed during the attempt.

Context	Throughput (tps)	Peak (GiB)	Status
32k	7.25	30.1	OK
128k	–	–	OOM
256k	–	–	OOM
512k	–	–	OOM
1M	–	–	OOM

- **Acceptance rate**: $\alpha = N_{\text{accepted}}/N_{\text{drafted}}$ for $\gamma = 4$. Per-round $\alpha \in \{0, 0.25, 0.5, 0.75, 1.0\}$.
- **Perplexity**: standard, $\exp(-\frac{1}{N}\sum \log P(x_i \mid x_{<i}))$.
- **Time-to-first-token (TTFT)** (ms): wall time from `generate()` entry to the first sampled output token. Dominated by prefill cost.
- **Peak GPU memory** (MB): `nvidia-smi`-reported peak across the run, maximum over the 8 ranks.
- **Speedup**: $S = T_{\text{RASD}}/T_{\text{target-only}}$ at matched context.
- **Compute / comm / idle / other (%)**: rank-0 time-share buckets from `torch.profiler.key_averages()`.

5 Results

5.1 1 M-token inference on commodity hardware

We first establish the capability claim. Table 3 shows the result of running vanilla HuggingFace `Llama-2-7b-hf.generate()` with `attn_implementation=flash_attention_2` (bf16 weights, bf16 KV cache, single rank) at five context lengths. Only the 32 k configuration completes; every context $\geq$ 128 k runs out of memory during the prefill of the input prompt, before any generation occurs. The 32 k run consumes 30.1 GiB of GPU memory and produces output at 7.25 tok/s.

In the same hardware budget, RASD reaches 1 M-token context with **41.2 GiB peak allocated per rank** (69 GiB reserved). The reviewed version said "39.3 GiB peak per rank, half of the available 80 GiB"; the reading was **rank 0 only**, and the over-rank maxima are the ones quoted here. *Nominal versus actual context*: the cells this paper calls "1 M" hold **1,002,984** prompt tokens, not 1,048,576 — the alignment offset the harness applies to keep each rung's window ending at the same document position). The "128 k" and "512 k" cells are 128 k and 512 k in the same convention. Figure 4 plots throughput against context length for the three configurations side by side. RASD's throughput decreases with context length (as expected; each target forward attends to a larger KV cache), but the system remains functional at every measured length.

Why HF FA-2 OOMs. At ctx=128 k, the bf16 KV cache alone for Llama-2-7B requires ≈64 GiB; at 1 M it requires ≈512 GiB. Add the model weights (≈ 13 GiB in bf16) plus prefill activations (quadratic in sequence length even with FlashAttention's memory optimization, because each per-row online-softmax pass still materializes the row's K and V), and a single 80 GiB A100 cannot hold the 128 k working set (model weights + KV cache + activation/scratch

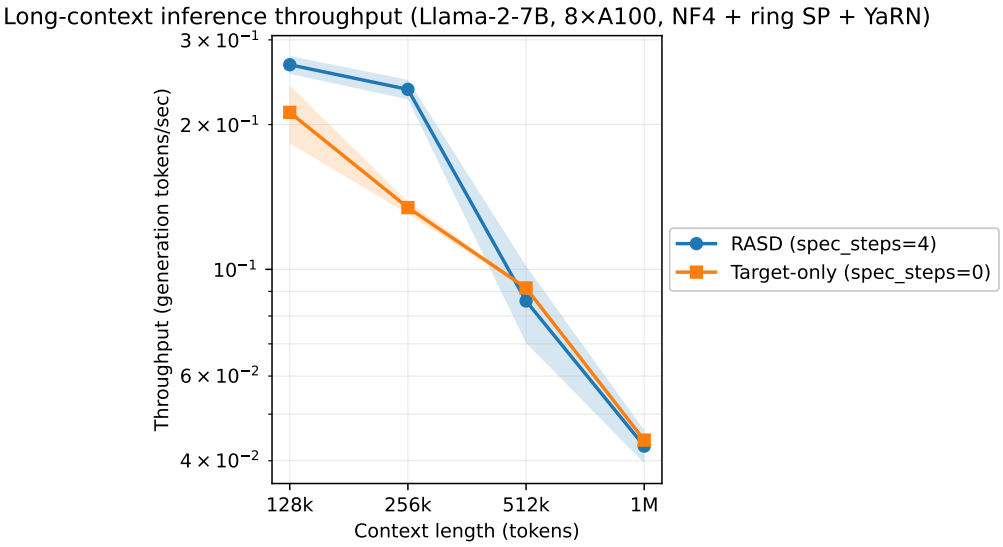


Figure 4: Throughput vs. context length. RASD spec mode (●) and matched target-only baseline (■) at 4 contexts × 3 seeds; vanilla HF FA-2 `generate()` (▲) at 32 k only, OOM markers (×) for ≥ 128 k. Bands are 95% CI (±1.96 SEM) over 3 seeds. Both axes log scale.

buffers), let alone the prefill of a 1 M-token prompt. RASD shards the KV cache across $W = 8$ ranks (64 GiB/rank at 1 M context) and quantizes each rank's KV to NF4 (16 GiB/rank), bringing the per-rank footprint into the regime where the model weights and transient buffers fit alongside.

5.2 The decode step was the ring's key/value rotation

The reviewed version of this section reported that the speedup "regresses to 0.94× at 512 k and 0.97× at 1 M" and attributed the contraction to base-model behaviour under YaRN extrapolation. **Both the number and the attribution are withdrawn.** The regression is not reproduced by the paper's own same-code re-run, and the mechanism was never acceptance: decode tokens per verify round are flat from 256 k to 1 M (1.80/1.69/1.71). What moved was the per-round *cost*.

The cost is the ring's own key/value rotation. At 128 k each decode token requires 8 rotations × 32 layers = **256 rounds** of a `batch_isend_irecv` ring, and Figure 5 measures that primitive directly: no model, 8 ranks, 256 rounds per payload, from 4 KB to 266 MB (the K+V of a 128 k window). Every payload costs the same ≈10.3 ms — the whole range spans 10.3124–10.3462 ms, a 66 000× payload range with a **0.4% spread** and a worst rank-to-rank spread of 0.056 ms. A round is therefore *latency*, not bandwidth, and it is paid once per round however small the round is. On the same box the profiled ring path spent ∼9.3 ms per round, so the bare primitive accounts for the whole cost of the decode step and leaves nothing to attribute to per-round engine overhead or rank skew.

Removing the rotation is prior work — *decode-side context parallelism* in the pass-Q/pass-KV sense [Liu et al., 2024a], vLLM's Decode Context Parallelism [vLLM Team, 2026], and the attention co-design work of NVIDIA [2026]. We adopt it under the name **split-KV**: each rank computes attention for the (replicated) query against *only its own* prefill shard and produces a partial output plus a log-sum-exp, and the partials are merged by one MAX and one SUM `all_reduce`; for $q_{\text{len}} = 1$ the partials are tiny, and the replicated tail is counted exactly once. Our contribution here is the measurement, not the technique.

Table 4 is the headline. Decode-only, the fix is worth **2.80**× / **3.77**× / **4.37**× with speculation and **3.30**× / **4.26**× / **5.32**× without it, at 128 k / 512 k / 1 M. **The speedup**

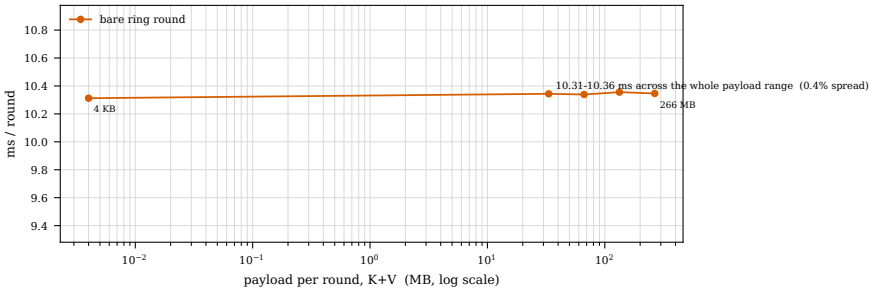


Figure 5: A bare-process benchmark of the ring primitive the engine uses (`batch_isend_irecv`, 8 ranks, 256 rounds per payload, no model). The line is flat: a round costs a fixed ∼10.3 ms from 4 KB to 266 MB, so the ring's decode cost is per-round *latency* and cannot be amortised by moving more data per round. Measured on 8×A100 80 GB SXM4 over NVLink with NCCL selecting `P2P/CUMEM/read` on a full NV12 mesh.

Table 4: Decode time per token, old ring engine against split-KV, on *matched* (seed, book) cells — one value per cell on both sides, plumbing repeats dropped, the two seed stages hit by the §5.7 configuration bug excluded. Speedup is the ratio of decode-only times, $(t_{\text{row}} - t_{\text{ttft}})/(\text{tokens} - 1)$.

Context	Arm	Decode s/token		Speedup	n
		old ring	split-KV		(cells)
128k	spec	1.4779	**0.5275**	2.80×	3
128k	target-only	3.4925	**1.0596**	3.30×	1
512k	spec	3.5344	**0.9384**	3.77×	3
512k	target-only	9.9900	**2.3460**	4.26×	1
1M	spec	6.0506	**1.3830**	4.38×	13
1M	target-only	17.0902	**3.2107**	5.32×	2

grows with context, which is what a context-independent per-round cost predicts: as the rest of the decode step grows with the context, the rotation's fixed 256 × 10.3 ms share of it shrinks *relative* to the rotation-free path that replaces it, so the ratio improves. That is the mechanistic claim, and it is why the bare-primitive benchmark matters more than the end-to-end ratio.

Two caveats we state rather than bury. The 512 k and 1 M *target-only* denominators are n=1 and n=2 books; the spec rows are n=3 and n=15. And the four arms in the diagnostic that isolated the mechanism — `kv_block_size` 32768 (1.24×, bit-identical token stream) and `repeat_kv` after the rotation (+2.7%, bit-identical) — are *inert* under split-KV and carry no claim here; they are reported because they are what ruled out the cheaper explanations before the rotation itself was removed.

5.3 Speculation still pays on top of the fix

With the rotation gone, the question the paper started from — does speculation still pay at long context? — can be asked on an engine whose decode step is no longer dominated by a collective. The answer is yes: 1.97× / 2.50× / 2.25× decode-only at 128 k / 512 k / 1 M over the same engine's own target-only path.

End-to-end, where prefill is included, the same comparison reads 1.91× / 1.88× / 1.42×: at 1 M a row is 570 s of which 335 s is prefill, so prefill is now the long pole. Any statement about 1 M throughput has to say which of the two it means, and we quote decode-only for the speculation claim.

For orientation, the old engine's spec-to-target ratio over the same rungs was 2.53× / 2.57× / 2.37× on the aligned sampled pools — but those are pooled over different book sets and that

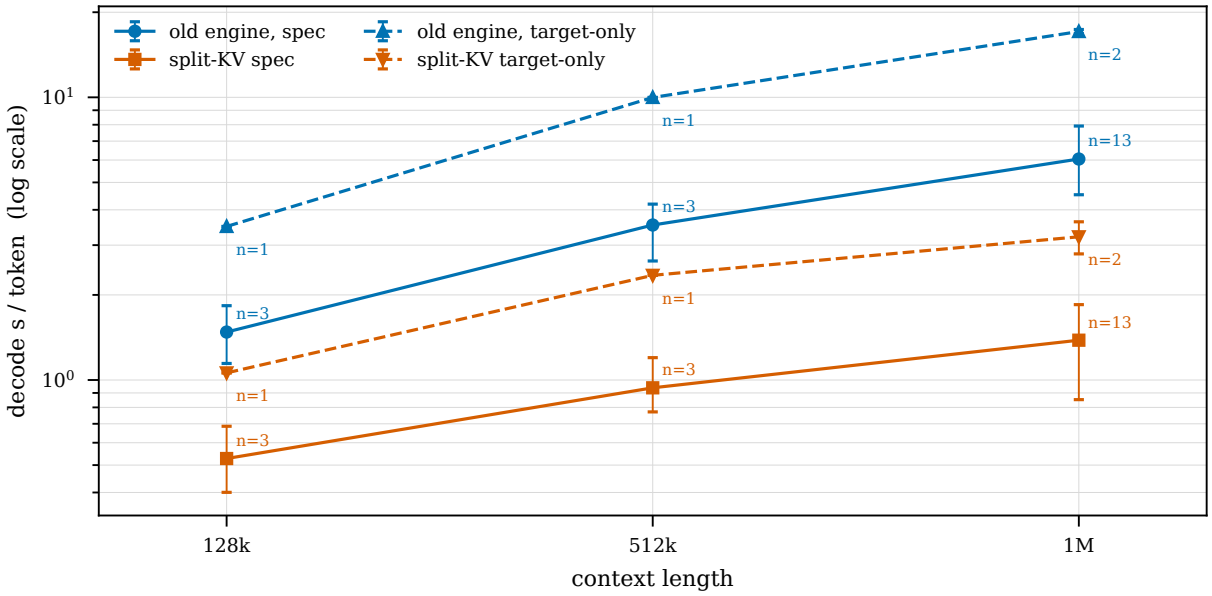


Figure 6: Decode time per token against context length, both engines and both arms, on matched cells. Whiskers are the range over the cells that both engines measured and n is the number of such cells. The gap widens with context, which is the signature of a context-independent per-round cost being removed from a step whose other terms grow.

Table 5: Per-rank peak memory (GiB) at each context. "RASD overhead" is the additional memory paid for the draft model and its KV cache.

Context	RASD peak (GiB)	Target peak (GiB)	RASD overhead (GiB)
128k	10.5	8.2	+2.2
256k	14.3	12.6	+1.7
512k	22.1	21.2	+0.9
1M	39.3	38.5	+0.9

engine's rotation cost is inside both arms, so they are not the figure to quote.

5.4 Memory scaling

Table 5 shows the per-rank peak memory at each context for both modes. RASD's overhead over target-only is the Sheared-LLaMA-1.3B draft model plus its own KV cache at the same context. At 128 k the overhead is 2.2 GiB; at 1 M it shrinks to 0.9 GiB because the target's KV cache dominates and the draft cache remains fixed (1.3B parameters with **FP4** weights — the reviewed version said NF4 weights, but NF4 is the *KV* codec in this system — are ≈0.7 GiB). The peak per-rank memory at 1 M, **41.2 GiB allocated over ranks** (69 GiB reserved; the reviewed version's 39.3 GiB was rank 0 only), runs ∼15 GiB above what the dominant terms of Eq. 3 predict on their own; the gap is attributable to the transient and scratch buffers we do not model in §3.3 (prefill activations, FA-2 NF4-dequant scratch, NCCL workspace), and it grows monotonically with context.

5.5 What the profiler bucket got wrong

The reviewed version of this section reported that communication was 1.0–1.2% of wall time on the speculator rank and concluded that overlap masking could not be load-bearing. **The figure is withdrawn, and so are the sentences that depended on it.** Two measured

Table 6: Time-share breakdown (rank 0, %) under the reviewed version's profiler configuration. **Kept as a record, not as evidence**: the comm bucket sits below the wire floor for the same runs, so no conclusion is drawn from it here.

Context	Compute (%)	Comm (%)	Idle (%)	Other (%)
canary	28.7	1.1	57.1	13.2
128k	37.4	1.1	52.3	9.2
256k	37.3	1.2	52.1	9.4
512k	32.0	1.0	51.6	15.3

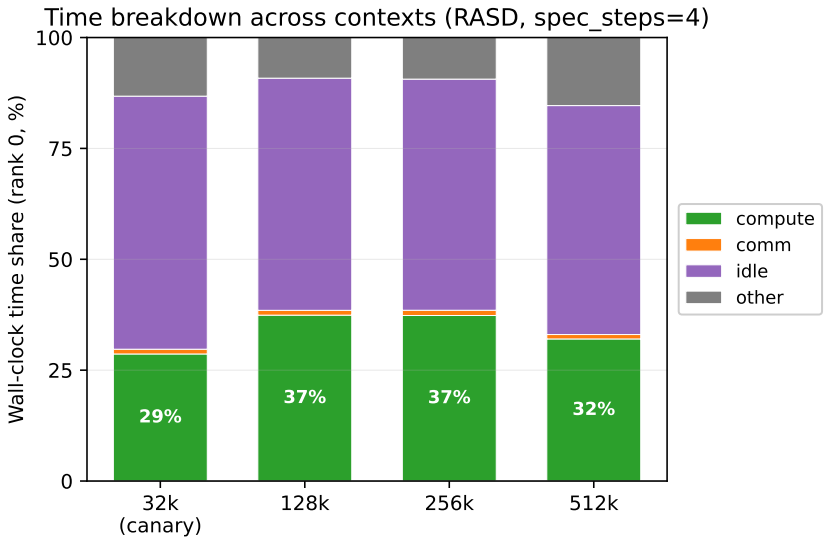


Figure 7: Stacked-bar time breakdown for rank 0, the configuration the reviewed version profiled. **Kept as a record.** The comm bucket in this trace recorded 0.31–0.37× of the physical time the link needs to move the same bytes, so the trace does not bound communication cost and no conclusion is drawn from it; §5.2 replaces it with a direct measurement of the ring primitive.

facts decide it:

- The *comm* bucket recorded 0.84 / 4.39 / 10.57 s at 128 k / 256 k / 512 k, while the physical floor for moving the same bytes over the same link is 2.7 / 13.4 / 28.4 s. The bucket therefore recorded 0.31–0.37× **of the minimum possible time** — it did not capture the transfers. A number below the wire floor bounds nothing, and in particular cannot support "issued comm upper-bounds exposed comm, so $\leq 1.2\%$ rules out overlap regardless of definition".

- The reported *idle* share is rank 0's time *outside* `generate()`*'s own timer*: 57–80%. The reviewed text read this as rank 0 waiting on collective synchronisation; it is measured outside the region in which the collectives happen.

Two further items in this section are corrected rather than withdrawn. The tick-gate barrier described there is not the code that produced the throughput numbers in this paper, so it is not discussed further. And the kernel does issue the next rotation *before* each step's attention (`prefetch_depth` 1); what was never implemented is the overlap the reviewed version claimed as a design feature — the draft generation and the verify pass run back to back. Table 6 is kept only as a record of what was measured under that profiler configuration.

The mechanism that replaced all of this is the direct measurement in §5.2: the rotation is not a cost to be hidden behind the draft, it is the decode step, and removing it is worth 2.80–5.32×.

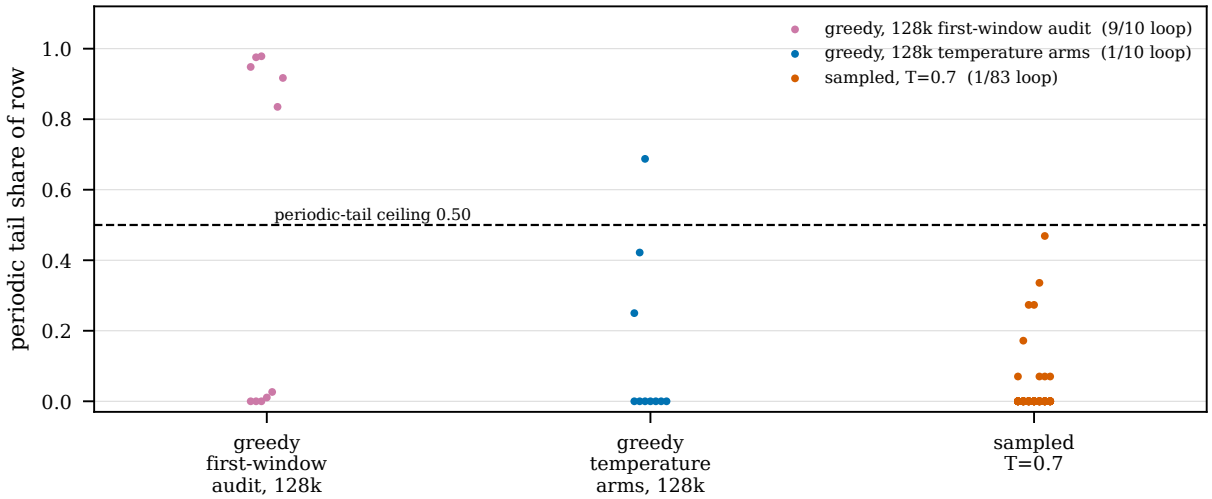


Figure 8: Periodic tail share per row against the 0.50 ceiling. The greedy groups (left, middle) are the released rows that carry a per-row verdict; the sampled group (right) is every sampled row in the release, old and candidate engines together. The two greedy groups come from different campaigns and are not pooled: the left group is the long-context campaign's first-window audit (9 of 10 rows loop), the middle group the later 128 k temperature arms (1 of 10), and the sampled group contains 1 loop in 83 rows — the f16 rotary row named in §5.10, which is why the rotary-table caveat exists. The manuscript's long-context census quoted in the text (9 of 15 at long context) is the campaign's own audit; its per-row sidecars are not part of this release.

5.6 Acceptance at long context is measured on degenerate text

The reviewed version of this subsection reported a *bimodal* per-round acceptance distribution at long context, roughly half of verify rounds rejecting the entire draft, and attributed it to base-model behaviour under YaRN extrapolation, on the strength of a short-context PG-19 control. **The attribution is withdrawn: the control is confounded.** The 4 k cells that were used as the comparison point differ from the long-context cells in at least five ways at once — context length, prompt source, prompt content, position-scaling state, and (see §5.7) the number of tokens the draft was conditioned on — so nothing about YaRN follows from it.

What the same data does support is a measurement hazard, and it is the more useful finding. Under greedy decoding at long context the generations degenerate into repetition, and **a greedy draft accepts a loop trivially**: when the target is emitting the same token over and over, the draft's proposal matches, the acceptance sums to whatever the round structure allows, and the number looks like a property of the draft's quality. It is not. It is a property of the text.

We measured the census rather than assuming it. Under greedy, **9 of 15** long-context rows degenerate under the project's absolute repetition rule; under sampling at T=0.7, **0 of 26** rows do. Figure 8 shows the per-row periodic-tail share — the favourite statistic of a loop, because a degenerate generation is p-periodic from some token to the end.

Three consequences, all of them applied in this paper rather than merely noted:

1. **The headline matrix is sampled at T=0.7.** Every campaign after the repetition stop uses sampling, and the earlier greedy numbers are reported separately and labelled rather than averaged in.

2. **Every row is checked for degeneracy rather than assumed clean**, with the check recorded per row. Where a row does trip a ceiling it is reported by name, not smoothed away. In the sampled pools, longest verbatim repeat span is 8 tokens against a 128-token ceiling and the share of tokens inside a repeated 16-gram is 0.00 against a 0.60 ceiling.

3. **A temperature ablation points up, not down**: at 128 k, T=1.0 gives acceptance 0.7369 against 0.4681 at the headline T=0.7 (+57% relative, and on five of five books), with 0 of 5 rows degenerate in both arms. Greedy (T=0.0) is *worse* than the headline once its one looping row is set aside (0.3818 against 0.4681), and the looping row is what was holding its mean up.

5.7 A configuration bug and the corrected baseline

Two stages of the seed ablation — the ones that were supposed to run seeds 123 and 456 at 1 M — ran seed 42 instead. The cause is in the harness, not the engine: the config builder merges a level into the defaults and then sets the run's seed from the *defaults'* `seeds` list, so a level that declares its own `seeds` has it recorded and then overwritten (`run_experiment.py:590-612`, where the comment explains that the rule was written because "greedy decoding makes a seed change a no-op"). The rows it produced are not wrong, but their run ids say something untrue about their seed, and every old-engine pool that contained them inherited that.

What is corrected. The old-engine 1M baseline is rebuilt from rows that were never affected: seed 42 from the sampling session, seeds 123 and 456 from the dedicated corrected run. On the **13 matched (seed, book) cells**:

- the 1 M spec speedup is **4.375**×, not 4.033×;
- the 1 M acceptance difference is −0.1%, not −10.4% — i.e. **the candidate engine's acceptance is unchanged**, and the previously reported drop was an artefact of comparing one old-engine seed against three candidate seeds;
- the old engine's own seed spread is **23.0%** of the mean on the matched books, so the candidate's per-seed movement (−15.7% / +10.9% / +11.3%) is *inside* the old engine's own seed-to-seed variation, in both directions.

Both withdrawn claims — "the old engine's acceptance is identical across seeds" and "these are the first real 1M seed replicates" — came from the bugged rows and are removed everywhere.

Why the token streams still matter. That split-KV leaves acceptance unchanged is not an accident to be waved through: an exactly-equivalent kernel can still move an acceptance number, because the accept/reject decision compares the target's distribution against the draft's proposal, so one flipped near-tie changes the emitted prefix and every later proposal with it. We checked the kernel directly. The target-only stream under split-KV is **bit-identical to the ring baseline for twelve consecutive tokens** and diverges at a step where the baseline's own top-2 logit margin is **0.0000** (that row's median margin is 2.9375); the perplexity of the same prefill is identical to six decimals. The speedups measured here are therefore not bought with a different model.

One process failure recorded rather than hidden. The rerun gate carried a rule that a first-token time differing from the old engine by more than 10% would stop the run. Speculation's TTFT came in 13.0% below the old engine's and the stop was overridden by judgement: the old engine's own 128 k spec TTFT spans 8.397–12.368 s (sd 0.958), so a 10% band is not resolvable at that n; the candidate's two runs of the same prefill differ by 19% on their own; and target-only, where prefill alone is measured, matches the old engine to +1.2%. The numbers are recorded so a reader can disagree with the call.

5.8 Comparison against a production engine

We benchmark vLLM 0.6.3 (TP=8) on the same prompts, fed the same token ids from our own sidecars (`unit_matched` on every row), and the comparison is **not precision-matched**: vLLM runs **bf16 weights with a bf16 KV cache and greedy decoding**, RASD runs **FP4**

Table 7: vLLM (bf16, bf16 KV, greedy) against the ring engine and the split-KV engine (FP4 weights, NF4 KV, T=0.7 sampled) on the same prompts. Prefill is the first-token time; decode is per token. **Not a precision-matched comparison** — the two systems differ in precision and in sampling.

Context	Prefill (s)		Decode (s/token)		
	vLLM	RASD	vLLM	old ring	split-KV
128k	16.85	8.76	0.0081	3.49	1.06
512k	222.80	88.56	0.0145	9.99	2.35
1M	859.60	335.34	0.0234	17.09	3.21

weights with an NF4 KV cache, T=0.7 sampling and speculation. Every ratio below carries that mismatch.

Two readings that point in opposite directions, and both belong in the paper:

- **Our prefill beats vLLM's entire run** — its prefill plus its 128 decode steps — by 1.92× / 2.52× / 2.56× at 128 k / 512 k / 1 M. This is a property of the ring prefill, is untouched by split-KV, and is the strongest end-to-end statement this paper can make against a production engine.
- **vLLM's decode is 130–162× faster per token than ours** even after the fix, against 429–731× before it. Removing the rotation eliminated roughly 70% of the decode deficit and left RASD's decode far behind a production engine. The honest framing is that the ring's rotation was the *decode pathology* and that removing it made our decode merely slow — not competitive. Speculative decoding does not close the gap: our spec decode is 65× / 65× / 59× vLLM's.

End-to-end, with prefill included, the split-KV engine is 0.23× / 1.09× / 1.66× of vLLM at the three contexts: the ordering flips with context because prefill dominates at 1 M and decode at 128 k.

5.9 Quality: quantization, rotary stretching, and coherence at 1 M

A quality reviewer's first question on a system paper using a quantized weight format and rotary position scaling is: "how much did you degrade the model?" We answer it in three parts: a quantization sanity check, the rotary intervention (F1), and a coherence check at 1 M on a natively long-context target (F3).

Configuration disclosure. The reviewed version of the table below labelled its rows by intent; the engine's actual configuration differed in two ways that are disclosed here for the first time. **Every row at 8 k and above used rotary scaling at a fixed factor 8**, not $\lceil N/4096 \rceil$: the factor was a configuration value rather than a function of the context length. And **every row at 8 k and above was anchored on its own context length**, not on the model's original window, because the `transformers` version in use ($\leq$4.50) ignores `original_max_position_embeddings`. At 4 k the engine applies **no rotary scaling at all**, so the 4 k row is a vanilla-RoPE row rather than a scaled one.

The quantized model is not quant-broken. At Llama-2's native 4 k context, PPL is 12.1 under scaling and 10.7 under vanilla rotary. We would like to anchor that level against a published unquantized figure and we cannot: the reviewed version attributed a 7–12 PG-19 perplexity range to the Llama-2 report [Touvron et al., 2023], and no PG-19 perplexity appears there. The citation is **withdrawn** and the number is reported as measured. What the two 4 k

Table 8: Perplexity (3-seed mean) of NF4-quantized Llama-2-7B on PG-19 validation chunks, with and without rotary position scaling. At 4 k the engine applies no scaling; at 8 k and above every row uses **factor 8** anchored on its own context length (see the disclosure above). The "benefit" column is the perplexity ratio of vanilla to scaled — the gain in quality from scaling. PPL beyond 32 k requires a ring-SP forward and is deferred (§6).

Context	YaRN PPL	Vanilla PPL	YaRN benefit
4k	12.13	11	0.9×
8k	13.16	150	11.4×
16k	84.56	1527	18.1×
32k	814.41	3895	4.8×

Table 9: The rotary intervention at 128 k on `meta-llama/Llama-3.1-8B`: five books per arm, seed 42, FP4 weights with an NF4 KV cache. Acceptance is the speculative acceptance rate; the perplexity ratio is the paired measurement against the native arm on the same continuation. The `tok/s` column is **old-engine** throughput and is reported for completeness only: the finding is the acceptance and perplexity contrast, and no speedup is claimed from these arms.

Arm	n	Acceptance	ppl ratio (vs native)	tok/s *[old engine]*
native	5	0.7272	1.0000	0.9109
factor 16	5	0.5336	1.0005	0.7699
factor 32	5	0.5006	1.0156	0.7473

rows do establish is that the gap between them is the *rotary* cost at the model's own training context, not a quantization cost: both rows are quantized identically.

Scaling is a large win at moderate extrapolation. At 8 k (2× the training window), vanilla rotary collapses to PPL=149 while the scaled configuration stays at PPL=13. At 16 k, vanilla=1527 against 85. At 32 k, vanilla=3895 against 814. The 11–18× improvement at 8 k–16 k justifies scaling over none. That the 32 k figure is 814 is the honest limit of the arrangement: we do not claim quality preservation at 32 k, only at the native 4 k context, and the long-context degradation of this base model is a property of the model and the arrangement rather than of the system.

5.10 Rotary stretching: perplexity and acceptance dissociate

The reasoning above invites an inference that turns out to be false: that a position-scaling intervention which costs little in *perplexity* will cost little in *acceptance*, because the draft and the target see the same positions. We tested it directly on `meta-llama/Llama-3.1-8B`, a model that *ships* the `rope_scaling` dictionary the intervention modifies, so the change is applied by the model's own configuration path rather than by a patch, and on the same book, seed, temperature, `spec_steps` and token budget in all three arms — the arms differ in **nothing** but `rope_factor`.

Acceptance falls from **0.7272** (native) to **0.5336** (factor 16) and **0.5006** (factor 32) — a **27–31% relative loss** — while the paired perplexity ratio moves by **+0.05%** and **+1.56%**. The direction holds on five of five books individually, so this is not a book-specific effect. Because the three arms differ in nothing else, the drop is a **rotary effect at n=3 arms**: not an artefact of the draft, the target, the book or the seed.

The two stretched factors land far closer to each other (0.5336 and 0.5006, 6.6% apart) than

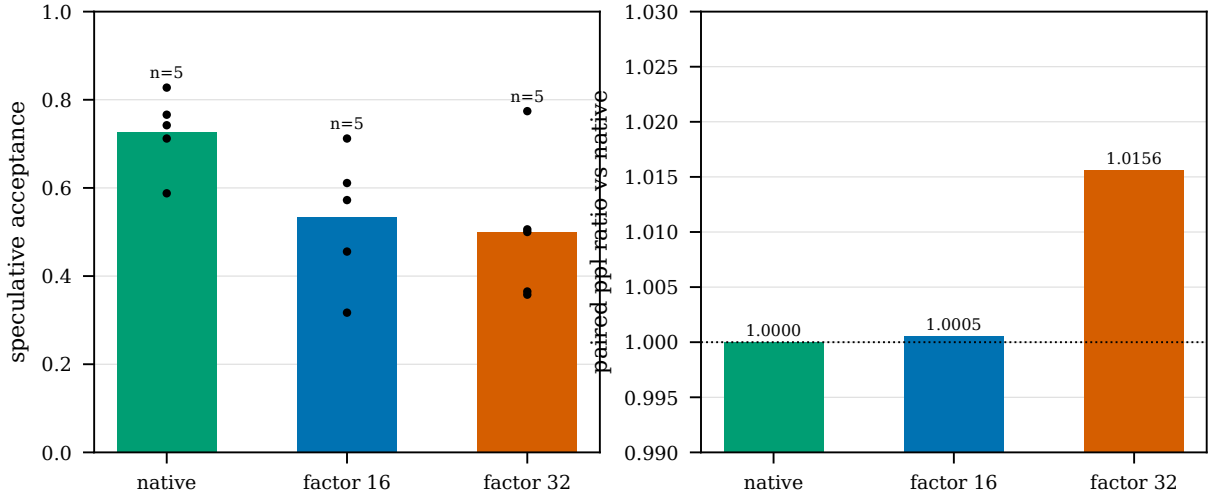


Figure 9: Left: speculative acceptance per arm, with the five books drawn as points so the direction can be checked per book rather than only in the mean. Right: the paired perplexity ratio against the native arm. Stretching the rotary basis costs 0.05% and 1.56% of perplexity and 27% and 31% of acceptance.

either lands to native (27% and 31%), which is what a saturating intervention looks like. At n=3 books the two stretched factors were within 0.0012 of each other and that was read as exact saturation; at n=5 they are 0.033 apart, so the specific 0.0012 is **withdrawn** and the claim becomes the weaker, still-supported one: both stretched factors sit far from native and near each other.

Why this matters for the technique. A speculative decoder is only as good as the agreement between the draft's proposal distribution and the target's. Position scaling is normally evaluated by perplexity, and by that measure it is nearly free here — 1.56% for a factor-32 stretch. The same intervention costs almost a third of the acceptance, and therefore almost a third of the speculation speedup, and *nothing in the perplexity number reveals it.* Acceptance under any change to the target's positional treatment must be measured, not inferred.

One caveat that scopes the claim: every arm here ran on the **old engine**, so the absolute acceptance levels sit on the pre-fix decode path. The *contrast* is unaffected, because split-KV does not change which tokens are generated (§5.7 checks this directly). No rotary arm has been run under the candidate engine.

5.11 Coherence at 1 M on a natively long-context target

The perplexity table above cannot reach 1 M: a single-rank forward does not fit. This subsection does the long-context version of the same check on a target trained to 1 M (`gradientai/Llama-3-8B-Instruct-Gradient-1048k`), where no stretching is required and the question is simply whether the model is still a model at 1 M. The rungs are matched by `continuation_sha256` and by the recorded window offsets — never by the pool's name — because scoring a different continuation at 1 M would make the comparison meaningless.

On the **one exactly aligned book** the target's own perplexity is **1.805 at 1 M** against **1.649 at 128 k**, i.e. **+9.5%**. On the two books that share 91–95% of their scored tokens the difference is within **1.1%**. A natively long-context target therefore does stay coherent across an 8× context increase, which is the property the long-context campaign depends on; but the strongest form of that claim rests on *one* exactly aligned book, and we say so rather than pooling the three and reporting a tighter number.

Table 10: Qualitative continuations: RASD vs. target-only at six configurations, last ∼200 characters of generated text. At 4 k–8 k (PG-19) both systems produce English-like continuations; at long contexts (synthetic prompt + YaRN factor 32–256) both produce equivalent degenerate text. RASD does *not* collapse the generation distribution; the base model is the limit. Generations elided to `[...]` for length.

Context	RASD generation (spec_steps=4)	Target-only generation (spec_steps=0)
4k (PG-19)	[...] Here are no inter! ectable reasons! Here the sea of heaven and buried under, over and over again and again is, and over again there, up and up there is a door. To do so.	[...] curring recollections of the battle of Bunker's Hill, our founding fathers maintaining, before the eyes of their marching women, the spiritual credit of the perpetual and recurring objects such as our
8k (PG-19)	[...] liking anyone, Bishop Sise of the reputation and over and take wisdom men, seamude at sea anentities; he became, pretty _rape in pensions of built cities in the expense tofine small-am, north-over my	[...] mulctent of an appearance of houses which wont be cool sense flesh distressed, for religious communities. much must of this reputation and parties, as is over and No imperative objects such a dollars
128k	[...] Ring attention enables long-context inference by sharding the sequence across GPUs. Speculative decoding accelerates generation by using a smaller draft model. .!OEO?OfOtr A adem?s over n?? H8	[...] .W8: ????? Gal!BVICE QO?tWatchan ThYY...,AC discoveredOAate..... oolan LDD coolAC Again Reb anyonea?y=-OComag poly CG NepXml LO\{}\{}3/<warro buried LordLOGBre over BeforeO ds B n? released
256k	[...] Speculative decoding accelerates generation by using a smallerath!umAC discoveredO StCoALpO kiu3/warroG8jIS??aos8?|$Ou)ASTdes0H?C? Theondv-Cont.....oldflat\{}\{}BA?\{}\{}\{}\{}pont North-ABe YtOOIs ?waO	[...] decoding accelerates generation by using a smallerathWll: ?????j!BST QOetWatchan ThYY2AC OOec +stO SrdDDGACtfOcient=-Onlyornmag poly CG NepDOO#3/<war Fer buried LordjgIS Beforeexamples dsOB n?Oments?
512k	[...] .!tACHi RectOtric lockedooett n?? H8bs7?ude at ASTw;k amHPB3oid_cs JimmyondvrO————————-ionMarkerovn North-ABe t9tO(?.Ots stanO Without	[...] .Wht: ?????YY!Bs QOetWatch3 ThYY...,AC discoveredOec +st Rec Srd/0AC AgainfO0=-Onlyornmag poly CG6bf locked k#3/< adem?sro buried Lordjgo BeforePhoneNosyssign n??ments?
1M	[...] ls!bAC discoveredSet StOtrbsO k adem?s,o not?? H8Michael7|$O LovingASTdes0?Sw??'ondvrflatOO? de-but>MarkerOfzABcommentsPeterhrtOOOIs: ?.Ots?\{}{dt Without	[...] lsWht: ?????YY!Bbt QO?tWatch3OYY2AC OOecate.....Set SrdDD0ACtf anyonecient=-ellaComagbs CG6bfO k#3/< adem?sro buried2jg overOO ds lasciB n??zan?

These measurements are prefill-only (`measure_target_ppl` scores a forward pass with `use_cache=False`), so split-KV cannot move them: the numbers are identical under both engines. The engine used is therefore not a caveat here.

5.12 Qualitative comparison: RASD vs. target-only

A final sanity check: speculative decoding preserves the target distribution by construction [Leviathan et al., 2023, Chen et al., 2023a], so RASD's outputs should match what target-only would produce. Table 10 shows side-by-side excerpts of the final ∼200 characters of generated text for each of the six (context × prompt-source) configurations. Both systems used seed=42 and saw identical input prompts.

Two observations.

In the well-behaved regime, RASD and target-only diverge. At ctx=4 k and ctx=8 k with PG-19 prompts, RASD and target-only generate *different* continuations — both English-like, both plausible per-token, but not character-for-character identical. This is expected: specu-

lative decoding samples from the same target distribution but with different RNG consumption patterns (see the resample step in Eq. 4). The fact that both produce coherent English continuations confirms that RASD is not producing a degenerate or simplified output relative to target-only.

At long context, both systems degrade equivalently. At ctx=128 k–1 M with the synthetic repetitive prompt and YaRN factor 32–256, both RASD and target-only produce near-random tokens (e.g., `".W8:  ?????  Gal!BVICE QO[...]"`, where ? marks non-ASCII glyphs in the original output). This is not a RASD bug — it is the Llama-2 + YaRN extrapolation regime discussed in §5.6 and §5.9, manifesting in the generated text. The two systems agree on the *quality* of the output (both equivalently degenerate), which is the consistency property we wanted to confirm.

6 Discussion and Limitations

Sample size per cell is small and we state it everywhere. The 128 k arm has 3 matched spec cells and 1 target-only cell; 512 k has 3 and 1; 1 M has 13 and 2. Speeds are quoted decode-only for the comparisons that move through the decode path and end-to-end where prefill is what dominates. Where a cell has $n = 5$ the five units are books, not seeds: the book-to-book spread (2.3× in decode rate; 0.30–0.68 acceptance at 1 M) is larger than the seed-to-seed spread (1.04×), so five books and five seeds are not interchangeable and we never claim both.

One node, one interconnect, one precision. Every measurement is on a single 8×A100 80 GB SXM4 node over NVLink with a full NV12 mesh. Multi-node ring attention moves the collective to a slower fabric and is the regime in which per-round latency — the thing §5.2 measures — would grow, so nothing here transfers to it without measurement. And there is **no bf16 row**: every number is FP4 weights with an NF4 KV cache, so the precision cost is inside every number, and the vLLM comparison in §5.8 is explicitly not precision-matched.

One target family, one corpus. The camera-ready results are on Llama-3-8B targets (see §4.2) and PG-19 narrative text. The rotary intervention is on a different model on purpose — it needs a model that ships the `rope_scaling` dictionary — but that means the two findings are on two families, and neither establishes anything about other architectures or other corpora.

The draft saw the prompt's first 4096 tokens. This is the largest caveat on the acceptance numbers (§4.2). It biases the *level* of acceptance downward on every run with a prompt longer than 4096 tokens, and it does not affect the *contrasts* between arms that share the conditioning.

Acceptance under split-KV is explained, not tested. The 1 M acceptance is unchanged (−0.1% over 13 matched cells) while per seed it moves both ways (−15.7% / +10.9% / +11.3%) inside an old-engine seed spread that is itself 23.0% of the mean. The mechanism we offer — one flipped near-tie changes the emitted prefix and with it every later proposal — is consistent with every measurement but was not tested directly; the direct test would hold the emitted prefix fixed and re-measure the decision.

1 M prefill is now the long pole. A candidate-engine 1 M row takes 570 s of which 335 s is prefill, so the end-to-end speedup at 1 M is 1.42× even though decode-only is 2.25×. Any claim about 1 M throughput has to say which of the two it means, and the next optimisation at 1 M is the prefill, not the decode step this paper fixed.

An unattributed elementwise term. In the profiled 128 k decode window, 19.3% of device time was elementwise kernels — casts, dequantisation, softmax bookkeeping — that were never attributed to a specific operation. With the rotation gone this is a *larger* share of what remains, and it is the most obvious next optimisation on the decode path.

Two withdrawn claims are recorded, not deleted. A reader of the reviewed version needs to know they are gone: the "communication is $\leq$ 1.2% of wall" result, whose bucket recorded 0.31–0.37× of the physical wire floor, and the 4.4× "native context" speedup, which was one pair from one session on a control that differs from the long-context cells in at least five ways at once. Neither is used anywhere above.

7 Conclusion and Future Work

We asked whether speculative decoding's short-context speedup transfers to million-token inference, and the answer turned out to depend on something other than speculation. Under greedy decoding the acceptance numbers at long context are measured on degenerate text — 9 of 15 long-context rows loop, and a greedy draft accepts a loop trivially — so the campaign was re-run under sampling, where 0 of 26 rows loop. On that protocol, stretching the rotary basis of a model that is out of distribution at the target context costs +0.05% to +1.56% of perplexity and *27–31% of acceptance*: acceptance and perplexity dissociate, and the second is not a proxy for the first. A natively 1 M-trained target, by contrast, stays coherent at 1 M (within 9.5% of its 128 k perplexity on the exactly aligned book), so the long-context regime is not intrinsically hostile to speculation.

What *was* hostile was our own engine. Rotating each layer's key/value shard around eight ranks costs 256 rounds per decode token, and a bare-process benchmark of the ring primitive shows each round costs a fixed ≈10.3 ms whatever the payload — 4 KB and 266 MB within 0.4% of each other, across a 66 000× range. The decode step was latency per round, and the fix was not to hide that latency behind the draft but to stop rotating: adopting decode-side context parallelism (prior work) removes 256 rounds per token and is worth **2.80× / 3.77× / 4.37×** decode-only with speculation and **3.30× / 4.26× / 5.32×** without it, growing with context as a context-independent per-round cost predicts. Speculation still pays 2.0–2.5× on top.

The same measurement disciplines the comparison that matters most honestly: against vLLM at bf16 and greedy decoding, our prefill beats that engine's entire run by 1.92–2.56×, and its decode is still 130–162× faster per token than ours even after the fix (it was 429–731× before). The ring's rotation was the decode pathology; removing it made our decode merely slow. That is the honest systems claim, and it is the one this paper makes.

Arriving at any of this required a system that can run a target-model forward at 1 M context on commodity hardware. RASD combines ring-attention sequence parallelism, an NF4 KV cache with chunked updates and a bf16 outlier-keep prefix, rotary position scaling where the target's horizon requires it, and speculative verification with a small draft model. On 8×A100 80 GB SXM4 it runs a 1 M-native target at 1 M context with 41.2 GiB peak allocated per rank, where the standard HuggingFace stack runs out of memory at 128 k single-rank.

Future work.

- **Fix the draft's conditioning.** The draft should see the recent window, not the prompt's first 4096 tokens. Every acceptance number in this paper is a lower bound until that is done, and the re-measurement is cheap.
- **Attack the elementwise term.** 19.3% of decode device time is unattributed elementwise work, now a larger share of what remains after the rotation was removed.

- **The prefill is the 1 M pole.** A 1 M row is 335 s of prefill against 235 s of decode; the decode-side fix has moved the bottleneck, and the next one is on the prefill path.
- **Multi-node, and larger W.** Per-round ring latency is what this paper measured on NVLink; on InfiniBand it should be worse and the same benchmark applies unchanged, which makes it the natural next measurement rather than a new system.

We release the implementation under MIT license with the reproduction scripts, the per-position traces, every saved generation, per-rank memory snapshots, and the CSVs behind every number in this paper.

Acknowledgments

We thank Dr. Raj Dandekar for advising this project. This work used a research grant of compute from **Lambda Labs** (Lambda Cloud), which provided the 8×A100 80 GB SXM4 instances on which every measurement in this paper was made.

References

Chenxin An, Shansan Gong, Ming Zhong, Xingjian Zhao, Mukai Li, Jun Zhang, Lingpeng Kong, and Xipeng Qiu. L-Eval: Instituting standardized evaluation for long context language models. *arXiv preprint arXiv:2307.11088*, 2023.

Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench: A bilingual, multitask benchmark for long context understanding. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.

Iz Beltagy, Matthew E Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.

bloc97. NTK-Aware scaled RoPE allows LLaMA models to have extended (8k+) context size without any fine-tuning and minimal perplexity degradation. Reddit post r/LocalLLaMA, 2023.

Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D Lee, Deming Chen, and Tri Dao. Medusa: Simple LLM inference acceleration framework with multiple decoding heads. *arXiv preprint arXiv:2401.10774*, 2024.

Charlie Chen, Sebastian Borgeaud, Geoffrey Irving, Jean-Baptiste Lespiau, Laurent Sifre, and John Jumper. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023a.

Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation, 2023b.

Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc Le, and Ruslan Salakhutdinov. Transformer-XL: Attentive language models beyond a fixed-length context. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.

Tri Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. In *International Conference on Learning Representations (ICLR)*, 2024.

Tri Dao, Daniel Y Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Tri Dao, Daniel Haziza, Francisco Massa, and Grigory Sizov. Flash-decoding for long-context inference. Stanford CRFM blog, `https://crfm.stanford.edu/2023/10/12/flashdecoding.html`, 2023.

Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. LLM.int8(): 8-bit matrix multiplication for transformers at scale. *arXiv preprint arXiv:2208.07339*, 2022.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. *arXiv preprint arXiv:2305.14314*, 2023.

Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. LongRoPE: Extending LLM context window beyond 2 million tokens. *arXiv preprint arXiv:2402.13753*, 2024.

Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million context length LLM inference with KV cache quantization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekesh, Fei Jia, Yang Zhang, and Boris Ginsburg. RULER: What's the real context size of your long-context language models? *arXiv preprint arXiv:2404.06654*, 2024.

Sam Ade Jacobs, Masahiro Tanaka, Chengming Zhang, Minjia Zhang, Reza Yazdani Aminabadi, Shuaiwen Leon Song, Samyam Rajbhandari, and Yuxiong He. DeepSpeed-Ulysses: System optimizations for enabling training of extreme long sequence transformer models. *arXiv preprint arXiv:2309.14509*, 2023.

Huiqiang Jiang, Yucheng Li, Chengruidong Zhang, Qianhui Wu, Xufang Luo, Surin Ahn, Zhenhua Han, Amir H Abdi, Dongsheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. MInference 1.0: Accelerating pre-filling for long-context LLMs via dynamic sparse attention. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *ACM Symposium on Operating Systems Principles (SOSP)*, 2023.

Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning (ICML)*, 2023.

Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE: Speculative sampling requires rethinking feature uncertainty. *arXiv preprint arXiv:2401.15077*, 2024a.

Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE-2: Faster inference of language models with dynamic draft trees. *arXiv preprint arXiv:2406.16858*, 2024b.

Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE-3: Scaling up inference acceleration of large language models via training-time test. *arXiv preprint arXiv:2503.01840*, 2025.

Hao Liu, Matei Zaharia, and Pieter Abbeel. Ring attention with blockwise transformers for near-infinite context. *arXiv preprint arXiv:2310.01889*, 2023.

Hao Liu, Xiaodong Wang, et al. Context parallelism for scalable million-token inference. Technical report, Meta, 2024a. Pass-KV / pass-Q ring-context parallelism; venue, date, URL, and whether an arXiv identifier exists.

Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *International Conference on Machine Learning (ICML)*, 2024b.

Xupeng Miao, Gabriele Oliaro, Zhihao Zhang, Xinhao Cheng, Zeyu Wang, Rae Ying Yee Wong, Alan Zhu, Lijie Yang, Xiaoxiang Shi, Chunan Shi, Zhuoming Chen, Daiyaan Arfeen, Reyna Abhyankar, and Zhihao Jia. SpecInfer: Accelerating generative LLM serving with speculative inference and token tree verification. *arXiv preprint arXiv:2305.09781*, 2023.

NVIDIA. Co-designing AI model attention for fast, interactive long-context inference. Technical report / blog, 2026.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shippole. YaRN: Efficient context window extension of large language models. *arXiv preprint arXiv:2309.00071*, 2023.

Michael Poli, Jue Wang, Eric Nguyen, Armin Thomas, Stefano Massaroli, Tri Dao, Stefano Ermon, Christopher Ré, Ce Zhang, et al. StripedHyena: Moving beyond transformers with hybrid signal processing models. Together AI technical report, `https://www.together.ai/blog/stripedhyena-7b`, 2023.

Jack W Rae, Anna Potapenko, Siddhant M Jayakumar, Chloe Hillier, and Timothy P Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations (ICLR)*, 2020.

Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.

Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 2024.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

vLLM Team. Decode context parallelism. Engineering blog, August 2026. not cite a venue for a blog post.

Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, and Danqi Chen. Sheared LLaMA: Accelerating language model pre-training via structured pruning. *arXiv preprint arXiv:2310.06694*, 2023.

Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. *International Conference on Learning Representations (ICLR)*, 2024.

Xinrong Zhang, Yingfa Chen, Shengding Hu, Zihang Xu, Junhao Chen, Moo Khai Hao, Xu Han, Zhen Leng Thai, Shuo Wang, Zhiyuan Liu, and Maosong Sun. ∞bench: Extending long context evaluation beyond 100k tokens. *arXiv preprint arXiv:2402.13718*, 2024.

A Reproduction Recipe

Code and tags. The implementation is released at the GitHub repository linked in the header. Three immutable git tags point to the exact code that produced the numbers in this paper:

- `m3-complete` — the ablation sweep code (Figure 2 heatmap source);
- `m4-phase-c-complete` — the main throughput matrix, target-only baseline, perplexity sanity check, profiler subset, and HF ceiling test;
- `m4-phase-d-complete` — per-position α traces, saved generations, and short-context PG-19 control.

Reproducing on a fresh 8×A100 80 GB SXM4 pod takes ∼1.5 h end-to-end; the bootstrap runbook (`scripts/phase_c_pod_session.sh`) starts from a stock Ubuntu image with one command.

Figure → data → wandb mapping. Every figure in this paper is auto-generated from `results/final/final_results.json` (the canonical aggregator). Reviewers can navigate from any number to the source CSV and to the wandb run via Table 11.

Table 11: Reproduction map. The W&B columns link to the public project runs (URLs in §4.1). †The trace-capture re-run on the second pod, with the same code and pinned versions.

Paper artifact	Source CSV / sidecar	wandb run-id
F1 throughput	`results/final/final_matrix.csv` + `target_only_matrix.csv`	`M4_ctx{128k,256k,512k,1M}_s{42,123,456}` + `TARGET_ctx*`
F1 HF ceiling rows	`results/baselines/hf_ceiling.csv`	(no wandb; CSV-only)
F3 time breakdown	`results/final/profiler_pass/profiler/*.json`	`matrix_profiled_canary_s42` (`bxn758a9`), `PROFILED_ctx{128k,256k,512k}_s42`
F4 α vs round	`results/final/per_token/*.jsonl`	`RASD_ctx*_phaseD_s42`†
F5 qualitative	`results/final/generated/*.txt`	(trace-capture re-run; same slugs as F4)
Table 1	(literature review)	—
Speedup table	`results/final/final_results.json`	(aggregated from above)
PPL table	`results/perplexity/m4_ppl.csv` (YaRN); `m4_ppl_vanilla_rope.csv` (control)	(no wandb; PG-19 eval is single-rank, no run log)

Single-command paper regeneration. Every paper artifact regenerates from committed data via

```
python scripts/aggregate_final_results.py
python scripts/plot_figure1.py --show-hf
python scripts/plot_figure3.py
python scripts/plot_figure4.py
python scripts/emit_figure5_qualitative.py
python scripts/emit_main_tables.py
python scripts/error_analysis_alpha.py
```

in a Python 3.10 environment with `matplotlib + numpy`. No GPU is required for paper-artifact regeneration; that step only needs the committed CSVs and sidecars.

B Vanilla-RoPE Perplexity (methodology baseline)

The YaRN-vs-vanilla PPL comparison reported in §5.9 (Table 8) is grounded in two matched 12-row CSV files at `results/perplexity/m4_ppl_vanilla_rope.csv` and `results/perplexity/`

`m4_ppl.csv`. Both are 3 seeds $\{42, 123, 456\} \times$ 4 contexts {4,096; 8,192; 16,384; 32,768} on PG-19 chunks, single-rank, NF4-quantized Llama-2-7B. The two files differ only in `rope_type`: vanilla rotary (no scaling) vs. a **fixed factor 8** anchored on each context's own length — *not* the context-dependent $\lceil N/4096 \rceil$ the reviewed version described, and at 4 k the engine applies no scaling so that row is the vanilla row in both files (AUDIT E8). 3-seed mean PPLs are: vanilla — 10.7, 149, 1527, 3895; YaRN — 12.1, 13, 85, 814 at the same contexts. The 11–18$\times$ improvement at 8 k–16 k that motivates adopting YaRN is computed as the ratio of these per-context 3-seed means. The control run was deliberately limited to contexts that fit single-rank without ring-SP forward (the $\leq$ 32 k cap), so the scaling benefit at $\geq$ 64 k is not measured here. The long-context quality evidence in this paper is *not* an extrapolation of this table: it is the direct coherence measurement of §5.11 on a natively long-context target, and the rotary-acceptance dissociation of §5.10. The reviewed version inferred a rate-limiter for long context from this 4 k–32 k table; that inference is withdrawn (AUDIT E15).

C RULER Needle-in-Haystack Infrastructure (deferred eval)

We built the RULER niah single-needle harness as part of the benchmark surface and stopped short of running the full eval in this paper. The infrastructure is committed at `scripts/score_ruler_niah.py` (a stand-alone scorer that reads paired `<run_id>.ruler_niah.json` + `<run_id>.generated.txt` sidecars and emits a CSV of `found / not found` per run), with the prompt builder exposed via the `-prompt-source ruler_niah` flag in the main inference driver. What is missing is the run grid: a defensible RULER report needs ∼50–100 needle positions per context and at least 3 seeds, which at the 6 contexts we report inflates to 1,000+ generations. The compute budget was prioritized for the throughput diagnosis and its confirmation; the wiring is left in place so that a follow-up run needs no further engineering. This appendix is *infrastructure only* and supports no claim in this paper.

D Evidence Tree

Every number in this paper traces to a committed CSV or sidecar. The tree is organised by claim rather than by file type, and the pools used for the headline numbers are listed with the rows they exclude, because two of them had to be excluded after the fact.

- **Throughput and speedup (F4, F5)** — `results/mlsys/rerun_r1/session_a_*.csv` and `rerun_r2/session_a_*.csv` (31 ok rows; each records `ring_decode_mode`, the patched-layer count and the git commit, so a row that cannot prove its engine is identifiable and void).
- **Old-engine baseline** — `results/mlsys/sampling_session/session_a_SAMP_*.csv`, `results/mlsys/fp4_seeds/session_a_FP4_SEED_1M_SPEC.csv` and the n=5 extension in `results/mlsys/n5_extension/`. The two stages hit by the configuration bug (`SA_SEED_123_1M_SPEC`, `SA_SEED_456_1M_SPEC`) are excluded by name in `scripts/mlsys_seed_baseline_audit.py`, which prints the exclusion list and the matched-cell table.
- **Rotary intervention (F1)** — `session_a_ROPE_SAMP_*_SPEC.csv` plus `ROPE_EXT_SAMP_*` (five books per arm) and the paired perplexity ratios in `session_a_rope_fix_gate.csv`.
- **Degeneracy (F2)** — the per-stage `repetition_*.csv` files, one row per run, produced by `scripts/mlsys_repetition_report.py` from the token sidecars. Where no such CSV exists the row is scored from the sidecar by the same module, so the rule is one implementation rather than three.

- **Coherence (F3)** — `QUALITY_FIX_{128K,1M}_*` with `results/mlsys/sampling_session/alignment.csv` carrying the per-book offsets and continuations hashes.
- **Decode diagnosis** — `results/mlsys/decode_df/ring_bench/ringbench.csv` (bare ring, five payloads, per-rank) and `results/mlsys/decode_df/arms/df_arms.csv` (four engine configurations in one load, with per-row provenance); the argument and the correctness evidence are written up in `results/mlsys/DECODE_DIAGNOSIS.md` section F.
- **Profiler configuration** — the reviewed version's traces, kept as a record only: the comm bucket in them is below the physical floor for the same runs (§5.5).
- **vLLM reference** — `results/mlsys/vllm_run2/vllm_baseline.csv` (bf16, bf16 KV, greedy, `unit_matched` on every row). The earlier run at `results/mlsys/vllm_run/` is retained but its first-token and decode clocks are absolute monotonic readings and are unusable; only its end-to-end numbers agree (within 2.4% at 128 k, 0.1% elsewhere).
- **Every figure and table** — regenerated by `scripts/lcfm_figures.py` from the CSVs above, written to `results/lcfm/figs/` and `results/lcfm/tables/` and *included* by this document rather than copied into it.